
PySAT Documentation

Release 0.1.7.dev22

Alexey Ignatiev, Joao Marques-Silva, Antonio Morgado

Jan 13, 2023

CONTENTS

1	API documentation	3
1.1	Core PySAT modules	3
1.1.1	Cardinality encodings (<code>pysat.card</code>)	3
1.1.2	Boolean formula manipulation (<code>pysat.formula</code>)	9
1.1.3	Pseudo-Boolean encodings (<code>pysat.pb</code>)	29
1.1.4	SAT solvers' API (<code>pysat.solvers</code>)	31
1.2	Supplementary examples package	44
1.2.1	Fu&Malik MaxSAT algorithm (<code>pysat.examples.fm</code>)	44
1.2.2	Hard formula generator (<code>pysat.examples.genhard</code>)	46
1.2.3	Minimum/minimal hitting set solver (<code>pysat.examples.hitman</code>)	49
1.2.4	LBX-like MCS enumerator (<code>pysat.examples.lbx</code>)	54
1.2.5	LSU algorithm for MaxSAT (<code>pysat.examples.lsu</code>)	57
1.2.6	CLD-like MCS enumerator (<code>pysat.examples.mcsls</code>)	60
1.2.7	An iterative model enumerator (<code>pysat.examples.models</code>)	63
1.2.8	A deletion-based MUS extractor (<code>pysat.examples.musx</code>)	64
1.2.9	OptUx optimal MUS enumerator (<code>pysat.examples.optux</code>)	66
1.2.10	RC2 MaxSAT solver (<code>pysat.examples.rc2</code>)	69
	Python Module Index	77
	Index	79

This site covers the usage and API documentation of the PySAT toolkit. For the basic information on what PySAT is, please, see [the main project website](#).

API DOCUMENTATION

The PySAT toolkit has four core modules: *card*, *formula*, *pb* and *solvers*. The three of them (*card*, *pb* and *solvers*) are Python wrappers for the code originally implemented in the C/C++ languages while the *formula* module is a *pure* Python module. Version *0.1.4.dev0* of PySAT brings a new module called *pb*, which is a wrapper for the basic functionality of a third-party library *PyPBLib* developed by the *Logic Optimization Group* of the University of Lleida.

1.1 Core PySAT modules

1.1.1 Cardinality encodings (*pysat.card*)

List of classes

<i>EncType</i>	This class represents a C-like enum type for choosing the cardinality encoding to use.
<i>CardEnc</i>	This abstract class is responsible for the creation of cardinality constraints encoded to a CNF formula.
<i>ITotalizer</i>	This class implements the iterative totalizer encoding ¹¹ .

Module description

This module provides access to various *cardinality constraint*¹ encodings to formulas in conjunctive normal form (CNF). These include pairwise², bitwise², ladder/regular³⁴, sequential counters⁵, sorting⁶ and cardinality networks⁷, totalizer⁸, modulo totalizer⁹, and modulo totalizer for *k*-cardinality¹⁰, as well as a *native* cardinality constraint representation supported by the *MiniCard* solver.

¹¹ Ruben Martins, Saurabh Joshi, Vasco M. Manquinho, Inês Lynce. *Incremental Cardinality Constraints for MaxSAT*. CP 2014. pp. 531-548

¹ Olivier Roussel, Vasco M. Manquinho. *Pseudo-Boolean and Cardinality Constraints*. Handbook of Satisfiability. 2009. pp. 695-733

² Steven David Prestwich. *CNF Encodings*. Handbook of Satisfiability. 2009. pp. 75-97

³ Carlos Ansótegui, Felip Manyà. *Mapping Problems with Finite-Domain Variables to Problems with Boolean Variables*. SAT (Selected Papers) 2004. pp. 1-15

⁴ Ian P. Gent, Peter Nightingale. *A New Encoding of AllDifferent Into SAT*. In International workshop on modelling and reformulating constraint satisfaction problems 2004. pp. 95-110

⁵ Carsten Sinz. *Towards an Optimal CNF Encoding of Boolean Cardinality Constraints*. CP 2005. pp. 827-831

⁶ Kenneth E. Batchner. *Sorting Networks and Their Applications*. AFIPS Spring Joint Computing Conference 1968. pp. 307-314

⁷ Roberto Asin, Robert Nieuwenhuis, Albert Oliveras, Enric Rodríguez-Carbonell. *Cardinality Networks and Their Applications*. SAT 2009. pp. 167-180

⁸ Olivier Bailleux, Yacine Bouffhad. *Efficient CNF Encoding of Boolean Cardinality Constraints*. CP 2003. pp. 108-122

⁹ Toru Ogawa, Yangyang Liu, Ryuzo Hasegawa, Miyuki Koshimura, Hiroshi Fujita. *Modulo Based CNF Encoding of Cardinality Constraints and Its Application to MaxSAT Solvers*. ICTAI 2013. pp. 9-17

¹⁰ António Morgado, Alexey Ignatiev, Joao Marques-Silva. *MSCG: Robust Core-Guided MaxSAT Solving*. System Description. JSAT 2015. vol. 9, pp. 129-134

A cardinality constraint is a constraint of the form: $\sum_{i=1}^n x_i \leq k$. Cardinality constraints are ubiquitous in practical problem formulations. Note that the implementation of the pairwise, bitwise, and ladder encodings can only deal with AtMost1 constraints, e.g. $\sum_{i=1}^n x_i \leq 1$.

Access to all cardinality encodings can be made through the main class of this module, which is [CardEnc](#).

Additionally, to the standard cardinality encodings that are basically “static” CNF formulas, the module is designed to be able to construct *incremental* cardinality encodings, i.e. those that can be incrementally extended at a later stage. At this point only the *iterative totalizer*⁷ encoding is supported. Iterative totalizer can be accessed with the use of the [ITotalizer](#) class.

Module details

class pysat.card.CardEnc

This abstract class is responsible for the creation of cardinality constraints encoded to a CNF formula. The class has three *class methods* for creating AtMostK, AtLeastK, and EqualsK constraints. Given a list of literals, an integer bound and an encoding type, each of these methods returns an object of class [pysat.formula.CNFPlus](#) representing the resulting CNF formula.

Since the class is abstract, there is no need to create an object of it. Instead, the methods should be called directly as class methods, e.g. `CardEnc.atmost(lits, bound)` or `CardEnc.equals(lits, bound)`. An example usage is the following:

```
>>> from pysat.card import *
>>> cnf = CardEnc.atmost(lits=[1, 2, 3], encoding=EncType.pairwise)
>>> print(cnf.clauses)
[[-1, -2], [-1, -3], [-2, -3]]
>>> cnf = CardEnc.equals(lits=[1, 2, 3], encoding=EncType.pairwise)
>>> print(cnf.clauses)
[[1, 2, 3], [-1, -2], [-1, -3], [-2, -3]]
```

classmethod atleast(lits, bound=1, top_id=None, vpool=None, encoding=1)

This method can be used for creating a CNF encoding of an AtLeastK constraint, i.e. of $\sum_{i=1}^n x_i \geq k$. The method takes 1 mandatory argument `lits` and 3 default arguments can be specified: `bound`, `top_id`, `vpool`, and `encoding`.

Parameters

- **lits** (*iterable(int)*) – a list of literals in the sum.
- **bound** (*int*) – the value of bound k .
- **top_id** (*integer or None*) – top variable identifier used so far.
- **vpool** ([IDPool](#)) – variable pool for counting the number of variables.
- **encoding** (*integer*) – identifier of the encoding to use.

Parameter `top_id` serves to increase integer identifiers of auxiliary variables introduced during the encoding process. This is helpful when augmenting an existing CNF formula with the new cardinality encoding to make sure there is no collision between identifiers of the variables. If specified, the identifiers of the first auxiliary variable will be `top_id+1`.

Instead of `top_id`, one may want to use a pool of variable identifiers `vpool`, which is automatically updated during the method call. In many circumstances, this is more convenient than using `top_id`. Also note that parameters `top_id` and `vpool` **cannot** be specified *simultaneously*.

The default value of `encoding` is `EncType.seqcounter`.

The method *translates* the `AtLeast` constraint into an `AtMost` constraint by *negating* the literals of `lits`, creating a new bound $n - k$ and invoking `CardEnc.atmost()` with the modified list of literals and the new bound.

Raises

- **CardEnc.NoSuchEncodingError** – if encoding does not exist.
- **ValueError** – if bound is meaningless for encoding.

Return type a `CNFPlus` object where the new clauses (or the new native atmost constraint) are stored.

classmethod `atmost(lits, bound=1, top_id=None, vpool=None, encoding=1)`

This method can be used for creating a CNF encoding of an `AtMostK` constraint, i.e. of $\sum_{i=1}^n x_i \leq k$. The method shares the arguments and the return type with method `CardEnc.atleast()`. Please, see it for details.

classmethod `equals(lits, bound=1, top_id=None, vpool=None, encoding=1)`

This method can be used for creating a CNF encoding of an `EqualsK` constraint, i.e. of $\sum_{i=1}^n x_i = k$. The method makes consecutive calls of both `CardEnc.atleast()` and `CardEnc.atmost()`. It shares the arguments and the return type with method `CardEnc.atleast()`. Please, see it for details.

class `pysat.card.EncType`

This class represents a C-like `enum` type for choosing the cardinality encoding to use. The values denoting the encodings are:

```
pairwise      = 0
seqcounter    = 1
sortnetwrk    = 2
cardnetwrk    = 3
bitwise       = 4
ladder        = 5
totalizer     = 6
mtotalizer    = 7
kmtotalizer   = 8
native        = 9
```

The desired encoding can be selected either directly by its integer identifier, e.g. 2, or by its alphabetical name, e.g. `EncType.sortnetwrk`.

Note that while most of the encodings are produced as a list of clauses, the “native” encoding of `MiniCard` is managed as one clause. Given an `AtMostK` constraint $\sum_{i=1}^n x_i \leq k$, the native encoding represents it as a pair `[lits, k]`, where `lits` is a list of size `n` containing literals in the sum.

class `pysat.card.ITotalizer(lits=[], ubound=1, top_id=None)`

This class implements the iterative totalizer encoding[?]. Note that `ITotalizer` can be used only for creating `AtMostK` constraints. In contrast to class `EncType`, this class is not abstract and its objects once created can be reused several times. The idea is that a *totalizer tree* can be extended, or the bound can be increased, as well as two totalizer trees can be merged into one.

The constructor of the class object takes 3 default arguments.

Parameters

- **lits** (`iterable(int)`) – a list of literals to sum.
- **ubound** (`int`) – the largest potential bound to use.
- **top_id** (`integer` or `None`) – top variable identifier used so far.

The encoding of the current tree can be accessed with the use of `CNF` variable stored as `self.cnf`. Potential bounds **are not** imposed by default but can be added as unit clauses in the final CNF formula. The bounds are stored in the list of Boolean variables as `self.rhs`. A concrete bound k can be enforced by considering a unit clause `-self.rhs[k]`. **Note** that `-self.rhs[0]` enforces all literals of the sum to be *false*.

An *ITotalizer* object should be deleted if it is not needed anymore.

Possible usage of the class is shown below:

```
>>> from pysat.card import ITotalizer
>>> t = ITotalizer(lits=[1, 2, 3], ubound=1)
>>> print(t.cnf.clauses)
[[-2, 4], [-1, 4], [-1, -2, 5], [-4, 6], [-5, 7], [-3, 6], [-3, -4, 7]]
>>> print(t.rhs)
[6, 7]
>>> t.delete()
```

Alternatively, an object can be created using the `with` keyword. In this case, the object is deleted automatically:

```
>>> from pysat.card import ITotalizer
>>> with ITotalizer(lits=[1, 2, 3], ubound=1) as t:
...     print(t.cnf.clauses)
[[-2, 4], [-1, 4], [-1, -2, 5], [-4, 6], [-5, 7], [-3, 6], [-3, -4, 7]]
...     print(t.rhs)
[6, 7]
```

`delete()`

Destroys a previously constructed *ITotalizer* object. Internal variables `self.cnf` and `self.rhs` get cleaned.

`extend(lits=[], ubound=None, top_id=None)`

Extends the list of literals in the sum and (if needed) increases a potential upper bound that can be imposed on the complete list of literals in the sum of an existing *ITotalizer* object to a new value.

Parameters

- `lits` (*iterable(int)*) – additional literals to be included in the sum.
- `ubound` (*int*) – a new upper bound.
- `top_id` (*integer or None*) – a new top variable identifier.

The top identifier `top_id` applied only if it is greater than the one used in `self`.

This method creates additional clauses encoding the existing totalizer tree augmented with new literals in the sum and updating the upper bound. As a result, it appends the new clauses to the list of clauses of *CNF* `self.cnf`. The number of newly created clauses is stored in variable `self.nof_new`.

Also, if the upper bound is updated, a list of bounds `self.rhs` gets increased and its length becomes `ubound+1`. Otherwise, it is updated with new values.

The method can be used in the following way:

```
>>> from pysat.card import ITotalizer
>>> t = ITotalizer(lits=[1, 2], ubound=1)
>>> print(t.cnf.clauses)
[[-2, 3], [-1, 3], [-1, -2, 4]]
>>> print(t.rhs)
[3, 4]
```

(continues on next page)

(continued from previous page)

```

>>>
>>> t.extend(lits=[5], ubound=2)
>>> print(t.cnf.clauses)
[[-2, 3], [-1, 3], [-1, -2, 4], [-5, 6], [-3, 6], [-4, 7], [-3, -5, 7], [-4, -5,
↪ 8]]
>>> print(t.cnf.clauses[-t.nof_new:])
[[-5, 6], [-3, 6], [-4, 7], [-3, -5, 7], [-4, -5, 8]]
>>> print(t.rhs)
[6, 7, 8]
>>> t.delete()

```

increase(ubound=1, top_id=None)

Increases a potential upper bound that can be imposed on the literals in the sum of an existing *ITotalizer* object to a new value.

Parameters

- **ubound** (*int*) – a new upper bound.
- **top_id** (*integer or None*) – a new top variable identifier.

The top identifier `top_id` applied only if it is greater than the one used in `self`.

This method creates additional clauses encoding the existing totalizer tree up to the new upper bound given and appends them to the list of clauses of *CNF* `self.cnf`. The number of newly created clauses is stored in variable `self.nof_new`.

Also, a list of bounds `self.rhs` gets increased and its length becomes `ubound+1`.

The method can be used in the following way:

```

>>> from pysat.card import ITotalizer
>>> t = ITotalizer(lits=[1, 2, 3], ubound=1)
>>> print(t.cnf.clauses)
[[-2, 4], [-1, 4], [-1, -2, 5], [-4, 6], [-5, 7], [-3, 6], [-3, -4, 7]]
>>> print(t.rhs)
[6, 7]
>>>
>>> t.increase(ubound=2)
>>> print(t.cnf.clauses)
[[-2, 4], [-1, 4], [-1, -2, 5], [-4, 6], [-5, 7], [-3, 6], [-3, -4, 7], [-3, -5,
↪ 8]]
>>> print(t.cnf.clauses[-t.nof_new:])
[[-3, -5, 8]]
>>> print(t.rhs)
[6, 7, 8]
>>> t.delete()

```

merge_with(another, ubound=None, top_id=None)

This method merges a tree of the current *ITotalizer* object, with a tree of another object and (if needed) increases a potential upper bound that can be imposed on the complete list of literals in the sum of an existing *ITotalizer* object to a new value.

Parameters

- **another** (*ITotalizer*) – another totalizer to merge with.

- **ubound** (*int*) – a new upper bound.
- **top_id** (*integer or None*) – a new top variable identifier.

The top identifier `top_id` applied only if it is greater than the one used in `self`.

This method creates additional clauses encoding the existing totalizer tree merged with another totalizer tree into *one* sum and updating the upper bound. As a result, it appends the new clauses to the list of clauses of *CNF* `self.cnf`. The number of newly created clauses is stored in variable `self.nof_new`.

Also, if the upper bound is updated, a list of bounds `self.rhs` gets increased and its length becomes `ubound+1`. Otherwise, it is updated with new values.

The method can be used in the following way:

```
>>> from pysat.card import ITotalizer
>>> with ITotalizer(lits=[1, 2], ubound=1) as t1:
...     print(t1.cnf.clauses)
[[-2, 3], [-1, 3], [-1, -2, 4]]
...     print(t1.rhs)
[3, 4]
...
...     t2 = ITotalizer(lits=[5, 6], ubound=1)
...     print(t1.cnf.clauses)
[[-6, 7], [-5, 7], [-5, -6, 8]]
...     print(t1.rhs)
[7, 8]
...
...     t1.merge_with(t2)
...     print(t1.cnf.clauses)
[[-2, 3], [-1, 3], [-1, -2, 4], [-6, 7], [-5, 7], [-5, -6, 8], [-7, 9], [-8, 10], [-3, 9], [-4, 10], [-3, -7, 10]]
...     print(t1.cnf.clauses[-t1.nof_new:])
[[-6, 7], [-5, 7], [-5, -6, 8], [-7, 9], [-8, 10], [-3, 9], [-4, 10], [-3, -7, 10]]
...     print(t1.rhs)
[9, 10]
...
...     t2.delete()
```

new (*lits=[]*, *ubound=1*, *top_id=None*)

The actual constructor of *ITotalizer*. Invoked from `self.__init__()`. Creates an object of *ITotalizer* given a list of literals in the sum, the largest potential bound to consider, as well as the top variable identifier used so far. See the description of *ITotalizer* for details.

exception `pysat.card.NoSuchEncodingError`

This exception is raised when creating an unknown an *AtMostK*, *AtLeastK*, or *EqualK* constraint encoding.

with_traceback ()

`Exception.with_traceback(tb)` – set `self.__traceback__` to `tb` and return `self`.

exception `pysat.card.UnsupportedBound`

This exception is raised when creating a pairwise, or bitwise, or ladder encoding of *AtMostK*, *AtLeastK*, or *EqualsK* with the bound different from 1 and $N - 1$.

with_traceback ()

`Exception.with_traceback(tb)` – set `self.__traceback__` to `tb` and return `self`.

1.1.2 Boolean formula manipulation (`pysat.formula`)

List of classes

<code>IDPool</code>	A simple manager of variable IDs.
<code>CNF</code>	Class for manipulating CNF formulas.
<code>CNFPlus</code>	CNF formulas augmented with <i>native</i> cardinality constraints.
<code>WCNF</code>	Class for manipulating partial (weighted) CNF formulas.
<code>WCNFPlus</code>	WCNF formulas augmented with <i>native</i> cardinality constraints.

Module description

This module is designed to facilitate fast and easy PySAT-development by providing a simple way to manipulate formulas in PySAT. Although only clausal formulas are supported at this point, future releases of PySAT are expected to implement data structures and methods to manipulate arbitrary Boolean formulas. The module implements the `CNF` class, which represents a formula in *conjunctive normal form* (CNF).

Recall that a CNF formula is conventionally seen as a set of clauses, each being a set of literals. A literal is a Boolean variable or its negation. In PySAT, a Boolean variable and a literal should be specified as an integer. For instance, a Boolean variable x_{25} is represented as integer 25. A literal $\neg x_{10}$ should be specified as -10. Moreover, a clause $(\neg x_2 \vee x_{19} \vee x_{46})$ should be specified as [-2, 19, 46] in PySAT. *Unit size clauses* are to be specified as unit size lists as well, e.g. a clause (x_3) is a list [3].

CNF formulas can be created as an object of class `CNF`. For instance, the following piece of code creates a CNF formula $(\neg x_1 \vee x_2) \wedge (\neg x_2 \vee x_3)$.

```
>>> from pysat.formula import CNF
>>> cnf = CNF()
>>> cnf.append([-1, 2])
>>> cnf.append([-2, 3])
```

The clauses of a formula can be accessed through the `clauses` variable of class `CNF`, which is a list of lists of integers:

```
>>> print(cnf.clauses)
[[-1, 2], [-2, 3]]
```

The number of variables in a CNF formula, i.e. the *largest variable identifier*, can be obtained using the `nv` variable, e.g.

```
>>> print(cnf.nv)
3
```

Class `CNF` has a few methods to read and write a CNF formula into a file or a string. The formula is read/written in the standard *DIMACS CNF* format. A clause in the DIMACS format is a string containing space-separated integer literals followed by 0. For instance, a clause $(\neg x_2 \vee x_{19} \vee x_{46})$ is written as -2 19 46 0 in DIMACS. The clauses in DIMACS should be preceded by a *preamble*, which is a line `p cnf nof_variables nof_clauses`, where `nof_variables` and `nof_clauses` are integers. A preamble line for formula $(\neg x_1 \vee x_2) \wedge (\neg x_2 \vee x_3)$ would be `p cnf 3 2`. The complete DIMACS file describing the formula looks this:

```
p cnf 3 2
-1 2 0
-2 3 0
```

Reading and writing formulas in DIMACS can be done with PySAT in the following way:

```
>>> from pysat.formula import CNF
>>> f1 = CNF(from_file='some-file-name.cnf') # reading from file
>>> f1.to_file('another-file-name.cnf') # writing to a file
>>>
>>> with open('some-file-name.cnf', 'r+') as fp:
...     f2 = CNF(from_fp=fp) # reading from a file pointer
...
...     fp.seek(0)
...     f2.to_fp(fp) # writing to a file pointer
>>>
>>> f3 = CNF(from_string='p cnf 3 3\n-1 2 0\n-2 3 0\n-3 0\n')
>>> print(f3.clauses)
[[-1, 2], [-2, 3], [-3]]
>>> print(f3.nv)
3
```

Besides plain CNF formulas, the `pysat.formula` module implements an additional class for dealing with *partial* and *weighted partial* CNF formulas, i.e. WCNF formulas. A WCNF formula is a conjunction of two sets of clauses: *hard* clauses and *soft* clauses, i.e. $\mathcal{F} = \mathcal{H} \wedge \mathcal{S}$. Soft clauses of a WCNF are labeled with integer *weights*, i.e. a soft clause of \mathcal{S} is a pair (c_i, w_i) . In partial (unweighted) formulas, all soft clauses have weight 1.

WCNF can be of help when solving optimization problems using the SAT technology. A typical example of where a WCNF formula can be used is *maximum satisfiability* (MaxSAT), which given a WCNF formula $\mathcal{F} = \mathcal{H} \wedge \mathcal{S}$ targets satisfying all its hard clauses \mathcal{H} and maximizing the sum of weights of satisfied soft clauses, i.e. maximizing the value of $\sum_{c_i \in \mathcal{S}} w_i \cdot c_i$.

An object of class `WCNF` has two variables to access the hard and soft clauses of the corresponding formula: `hard` and `soft`. The weights of soft clauses are stored in variable `wght`.

```
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF()
>>> wcnf.append([-1, -2])
>>> wcnf.append([1], weight=1)
>>> wcnf.append([2], weight=3) # the formula becomes unsatisfiable
>>>
>>> print(wcnf.hard)
[[-1, -2]]
>>> print(wcnf.soft)
[[1], [2]]
>>> print(wcnf.wght)
[1, 3]
```

A properly constructed WCNF formula must have a *top weight*, which should be equal to $1 + \sum_{c_i \in \mathcal{S}} w_i$. Top weight of a formula can be accessed through variable `topw`.

```
>>> wcnf.topw = sum(wcnf.wght) + 1 # (1 + 3) + 1
>>> print(wcnf.topw)
5
```

Note: Although it is not aligned with the WCNF format description, starting with the 0.1.5.dev8 release, PySAT is able to deal with WCNF formulas having not only integer clause weights but also weights represented as *floating point numbers*. Moreover, *negative weights* are also supported.

Additionally to classes `CNF` and `WCNF`, the module provides the extended classes `CNFPlus` and `WCNFPlus`. The only difference between `?CNF` and `?CNFPlus` is the support for *native* cardinality constraints provided by the `MiniCard` solver (see `pysat.card` for details). The corresponding variable in objects of `CNFPlus` (`WCNFPlus`, resp.) responsible for storing the `AtMostK` constraints is `atmosts` (`atms`, resp.). **Note** that at this point, `AtMostK` constraints in `WCNF` can be *hard* only.

Besides the implementations of `CNF` and `WCNF` formulas in PySAT, the `pysat.formula` module also provides a way to manage variable identifiers. This can be done with the use of the `IDPool` manager. With the use of the `CNF` and `WCNF` classes as well as with the `IDPool` variable manager, it is pretty easy to develop practical problem encoders into SAT or MaxSAT/MinSAT. As an example, a PHP formula encoder is shown below (the implementation can also be found in `examples.genhard.PHP`).

```
from pysat.formula import CNF
cnf = CNF() # we will store the formula here

# nof_holes is given

# initializing the pool of variable ids
vpool = IDPool(start_from=1)
pigeon = lambda i, j: vpool.id('pigeon{0}@{1}'.format(i, j))

# placing all pigeons into holes
for i in range(1, nof_holes + 2):
    cnf.append([pigeon(i, j) for j in range(1, nof_holes + 1)])

# there cannot be more than 1 pigeon in a hole
pigeons = range(1, nof_holes + 2)
for j in range(1, nof_holes + 1):
    for comb in itertools.combinations(pigeons, 2):
        cnf.append([-pigeon(i, j) for i in comb])
```

Module details

```
class pysat.formula.CNF(from_file=None, from_fp=None, from_string=None, from_clauses=[],
                        from_aiger=None, comment_lead=['c'])
```

Class for manipulating CNF formulas. It can be used for creating formulas, reading them from a file, or writing them to a file. The `comment_lead` parameter can be helpful when one needs to parse specific comment lines starting not with character `c` but with another character or a string.

Parameters

- **from_file** (*str*) – a DIMACS CNF filename to read from
- **from_fp** (*file_pointer*) – a file pointer to read from
- **from_string** (*str*) – a string storing a CNF formula
- **from_clauses** (*list(list(int))*) – a list of clauses to bootstrap the formula with
- **from_aiger** (*aiger.AIG* (see `py-aiger` package)) – an AIGER circuit to bootstrap the formula with

- **comment_lead** (*list(str)*) – a list of characters leading comment lines

append(*clause*)

Add one more clause to CNF formula. This method additionally updates the number of variables, i.e. variable `self.nv`, used in the formula.

Parameters **clause** (*list(int)*) – a new clause to add.

```
>>> from pysat.formula import CNF
>>> cnf = CNF(from_clauses=[[-1, 2], [3]])
>>> cnf.append([-3, 4])
>>> print(cnf.clauses)
[[-1, 2], [3], [-3, 4]]
```

copy()

This method can be used for creating a copy of a CNF object. It creates another object of the `CNF` class and makes use of the *deepcopy* functionality to copy the clauses.

Returns an object of class `CNF`.

Example:

```
>>> cnf1 = CNF(from_clauses=[[-1, 2], [1]])
>>> cnf2 = cnf1.copy()
>>> print(cnf2.clauses)
[[-1, 2], [1]]
>>> print(cnf2.nv)
2
```

extend(*clauses*)

Add several clauses to CNF formula. The clauses should be given in the form of list. For every clause in the list, method `append()` is invoked.

Parameters **clauses** (*list(list(int))*) – a list of new clauses to add.

Example:

```
>>> from pysat.formula import CNF
>>> cnf = CNF(from_clauses=[[-1, 2], [3]])
>>> cnf.extend([[-3, 4], [5, 6]])
>>> print(cnf.clauses)
[[-1, 2], [3], [-3, 4], [5, 6]]
```

from_aiger(*aig*, *vpool=None*)

Create a CNF formula by Tseitin-encoding an input AIGER circuit.

Input circuit is expected to be an object of class `aiger.AIG`. Alternatively, it can be specified as an `aiger.BoolExpr`, or an `*.aag` filename, or an AIGER string to parse. (Classes `aiger.AIG` and `aiger.BoolExpr` are defined in the *py-aiger* package.)

Parameters

- **aig** (`aiger.AIG` (see *py-aiger* package)) – an input AIGER circuit
- **vpool** (`IDPool`) – pool of variable identifiers (optional)

Example:


```

>>> import aiger
>>> x, y, z = aiger.atom('x'), aiger.atom('y'), aiger.atom('z')
>>> expr = ~(x | y) & z
>>> print(expr.aig)
aag 5 3 0 1 2
2
4
8
10
6 3 5
10 6 8
i0 y
i1 x
i2 z
o0 6c454aea-c9e1-11e9-bbe3-3af9d34370a9
>>>
>>> from pysat.formula import CNF
>>> cnf = CNF(from_aiger=expr.aig)
>>> print(cnf.nv)
5
>>> print(cnf.clauses)
[[3, 2, 4], [-3, -4], [-2, -4], [-4, -1, 5], [4, -5], [1, -5]]
>>> print(['{0} <=> {1}'.format(v, cnf.vpool.obj(v)) for v in cnf.inps])
['3 <=> y', '2 <=> x', '1 <=> z']
>>> print(['{0} <=> {1}'.format(v, cnf.vpool.obj(v)) for v in cnf.outs])
['5 <=> 6c454aea-c9e1-11e9-bbe3-3af9d34370a9']

```

from_clauses(*clauses*)

This methods copies a list of clauses into a CNF object.

Parameters *clauses* (*list(list(int))*) – a list of clauses

Example:

```

>>> from pysat.formula import CNF
>>> cnf = CNF(from_clauses=[[-1, 2], [1, -2], [5]])
>>> print(cnf.clauses)
[[-1, 2], [1, -2], [5]]
>>> print(cnf.nv)
5

```

from_file(*fname*, *comment_lead*='c', *compressed_with*='use_ext')

Read a CNF formula from a file in the DIMACS format. A file name is expected as an argument. A default argument is *comment_lead* for parsing comment lines. A given file can be compressed by either gzip, bzip2, or lzma.

Parameters

- **fname** (*str*) – name of a file to parse.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines
- **compressed_with** (*str*) – file compression algorithm

Note that the *compressed_with* parameter can be None (i.e. the file is uncompressed), 'gzip', 'bzip2', 'lzma', or 'use_ext'. The latter value indicates that compression type should be automatically deter-

mined based on the file extension. Using 'lzma' in Python 2 requires the `backports.lzma` package to be additionally installed.

Usage example:

```
>>> from pysat.formula import CNF
>>> cnf1 = CNF()
>>> cnf1.from_file('some-file.cnf.gz', compressed_with='gzip')
>>>
>>> cnf2 = CNF(from_file='another-file.cnf')
```

from_fp(*file_pointer*, *comment_lead*=['c'])

Read a CNF formula from a file pointer. A file pointer should be specified as an argument. The only default argument is `comment_lead`, which can be used for parsing specific comment lines.

Parameters

- **file_pointer** (*file pointer*) – a file pointer to read the formula from.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines

Usage example:

```
>>> with open('some-file.cnf', 'r') as fp:
...     cnf1 = CNF()
...     cnf1.from_fp(fp)
>>>
>>> with open('another-file.cnf', 'r') as fp:
...     cnf2 = CNF(from_fp=fp)
```

from_string(*string*, *comment_lead*=['c'])

Read a CNF formula from a string. The string should be specified as an argument and should be in the DIMACS CNF format. The only default argument is `comment_lead`, which can be used for parsing specific comment lines.

Parameters

- **string** (*str*) – a string containing the formula in DIMACS.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines

Example:

```
>>> from pysat.formula import CNF
>>> cnf1 = CNF()
>>> cnf1.from_string('p cnf 2 2\n-1 2 0\n1 -2 0')
>>> print(cnf1.clauses)
[[-1, 2], [1, -2]]
>>>
>>> cnf2 = CNF(from_string='p cnf 3 3\n-1 2 0\n-2 3 0\n-3 0\n')
>>> print(cnf2.clauses)
[[-1, 2], [-2, 3], [-3]]
>>> print(cnf2.nv)
3
```

negate(*topv=None*)

Given a CNF formula \mathcal{F} , this method creates a CNF formula $\neg\mathcal{F}$. The negation of the formula is encoded

to CNF with the use of *auxiliary* Tseitin variables¹. A new CNF formula is returned keeping all the newly introduced variables that can be accessed through the `auxvars` variable.

Note that the negation of each clause is encoded with one auxiliary variable if it is not unit size. Otherwise, no auxiliary variable is introduced.

Parameters `topv` (*int*) – top variable identifier if any.

Returns an object of class `CNF`.

```
>>> from pysat.formula import CNF
>>> pos = CNF(from_clauses=[[-1, 2], [3]])
>>> neg = pos.negate()
>>> print(neg.clauses)
[[1, -4], [-2, -4], [-1, 2, 4], [4, -3]]
>>> print(neg.auxvars)
[4, -3]
```

to_alien(*file_pointer*, *format*='opb', *comments*=None)

The method can be used to dump a CNF formula into a file pointer in an alien file format, which at this point can either be LP, OPB, or SMT. The file pointer is expected as an argument. Additionally, the target format 'lp', 'opb', or 'smt' may be specified (equal to 'opb' by default). Finally, supplementary comment lines can be specified in the `comments` parameter.

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **format** (*str*) – alien file format to use
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import CNF
>>> cnf = CNF()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.lp', 'w') as fp:
...     cnf.to_alien(fp, format='lp') # writing to the file pointer
```

to_file(*fname*, *comments*=None, *compress_with*='use_ext')

The method is for saving a CNF formula into a file in the DIMACS CNF format. A file name is expected as an argument. Additionally, supplementary comment lines can be specified in the `comments` parameter. Also, a file can be compressed using either gzip, bzip2, or lzma (xz).

Parameters

- **fname** (*str*) – a file name where to store the formula.
- **comments** (*list(str)*) – additional comments to put in the file.
- **compress_with** (*str*) – file compression algorithm

Note that the `compress_with` parameter can be None (i.e. the file is uncompressed), 'gzip', 'bzip2', 'lzma', or 'use_ext'. The latter value indicates that compression type should be automatically determined based on the file extension. Using 'lzma' in Python 2 requires the `backports.lzma` package to be additionally installed.

¹ G. S. Tseitin. *On the complexity of derivations in the propositional calculus*. Studies in Mathematics and Mathematical Logic, Part II. pp. 115–125, 1968

Example:

```
>>> from pysat.formula import CNF
>>> cnf = CNF()
...
>>> # the formula is filled with a bunch of clauses
>>> cnf.to_file('some-file-name.cnf') # writing to a file
```

to_fp(file_pointer, comments=None)

The method can be used to save a CNF formula into a file pointer. The file pointer is expected as an argument. Additionally, supplementary comment lines can be specified in the `comments` parameter.

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import CNF
>>> cnf = CNF()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.cnf', 'w') as fp:
...     cnf.to_fp(fp) # writing to the file pointer
```

weighted()

This method creates a weighted copy of the internal formula. As a result, an object of class *WCNF* is returned. Every clause of the CNF formula is *soft* in the new *WCNF* formula and its weight is equal to 1. The set of hard clauses of the formula is empty.

Returns an object of class *WCNF*.

Example:

```
>>> from pysat.formula import CNF
>>> cnf = CNF(from_clauses=[[-1, 2], [3, 4]])
>>>
>>> wcnf = cnf.weighted()
>>> print(wcnf.hard)
[]
>>> print(wcnf.soft)
[[-1, 2], [3, 4]]
>>> print(wcnf.wght)
[1, 1]
```

class pysat.formula.**CNFPlus**(from_file=None, from_fp=None, from_string=None, comment_lead=['c'])

CNF formulas augmented with *native* cardinality constraints.

This class inherits most of the functionality of the *CNF* class. The only difference between the two is that *CNFPlus* supports *native* cardinality constraints of *MiniCard*.

The parser of input DIMACS files of *CNFPlus* assumes the syntax of *AtMostK* and *AtLeastK* constraints defined in the [description](#) of *MiniCard*:

```
c Example: Two cardinality constraints followed by a clause
p cnf+ 7 3
1 -2 3 5 -7 <= 3
4 5 6 -7 >= 2
3 5 7 0
```

Each `AtLeastK` constraint is translated into an `AtMostK` constraint in the standard way: $\sum_{i=1}^n x_i \geq k \leftrightarrow \sum_{i=1}^n \neg x_i \leq (n - k)$. Internally, `AtMostK` constraints are stored in variable `atmosts`, each being a pair `(lits, k)`, where `lits` is a list of literals in the sum and `k` is the upper bound.

Example:

```
>>> from pysat.formula import CNFPlus
>>> cnf = CNFPlus(from_string='p cnf+ 7 3\n1 -2 3 5 -7 <= 3\n4 5 6 -7 >= 2\n3 5 7 0\n')
>>> print(cnf.clauses)
[[3, 5, 7]]
>>> print(cnf.atmosts)
[[[1, -2, 3, 5, -7], 3], [[-4, -5, -6, 7], 2]]
>>> print(cnf.nv)
7
```

For details on the functionality, see [CNF](#).

append(*clause*, *is_atmost=False*)

Add a single clause or a single `AtMostK` constraint to `CNF+` formula. This method additionally updates the number of variables, i.e. variable `self.nv`, used in the formula.

If the clause is an `AtMostK` constraint, this should be set with the use of the additional default argument `is_atmost`, which is set to `False` by default.

Parameters

- **clause** (*list(int)*) – a new clause to add.
- **is_atmost** (*bool*) – if `True`, the clause is `AtMostK`.

```
>>> from pysat.formula import CNFPlus
>>> cnf = CNFPlus()
>>> cnf.append([-3, 4])
>>> cnf.append([1, 2, 3], 1, is_atmost=True)
>>> print(cnf.clauses)
[[-3, 4]]
>>> print(cnf.atmosts)
[[1, 2, 3], 1]
```

copy()

This method can be used for creating a copy of a `CNFPlus` object. It creates another object of the `CNFPlus` class, call the copy function of `CNF` class and makes use of the *deepcopy* functionality to copy the `atmost` constraints.

Returns an object of class `CNFPlus`.

Example:

```
>>> cnf1 = CNFPlus()
>>> cnf1.extend([-1, 2], [1])
```

(continues on next page)

(continued from previous page)

```

>>> cnf1.append([[1, 2], 1], is_atmost=True)
>>> cnf2 = cnf1.copy()
>>> print(cnf2.clauses)
[[-1, 2], [1]]
>>> print(cnf2.nv)
2
>>> print(cnf2.atmosts)
[[[1, 2], 1]]

```

extend(*formula*)

Extend the CNF+ formula with more clauses and/or AtMostK constraints. The additional clauses and AtMostK constraints to add should be given in the form of *CNFPlus*. Alternatively, a list of clauses can be added too. For every single clause and AtMostK constraint in the input formula, method *append()* is invoked.

Parameters *formula* (*CNFPlus*) – new constraints to add.

Example:

```

>>> from pysat.formula import CNFPlus
>>> cnf1 = CNFPlus()
>>> cnf1.extend([[ -3, 4], [5, 6], [[1, 2, 3], 1]])
>>> print(cnf1.clauses)
[[-3, 4], [5, 6]]
>>> print(cnf1.atmosts)
[[[1, 2, 3], 1]]
>>> cnf2 = CNFPlus()
>>> cnf2.extend(cnf1)
>>> print(cnf1.clauses)
[[-3, 4], [5, 6]]
>>> print(cnf1.atmosts)
[[[1, 2, 3], 1]]

```

from_fp(*file_pointer*, *comment_lead*=['c'])

Read a CNF+ formula from a file pointer. A file pointer should be specified as an argument. The only default argument is *comment_lead*, which can be used for parsing specific comment lines.

Parameters

- **file_pointer** (*file pointer*) – a file pointer to read the formula from.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines

Usage example:

```

>>> with open('some-file.cnf+', 'r') as fp:
...     cnf1 = CNFPlus()
...     cnf1.from_fp(fp)
>>>
>>> with open('another-file.cnf+', 'r') as fp:
...     cnf2 = CNFPlus(from_fp=fp)

```

to_alien(*file_pointer*, *format*='opb', *comments*=None)

The method can be used to dump a CNF+ formula into a file pointer in an alien file format, which at this point can either be LP, OPB, or SMT. The file pointer is expected as an argument. Additionally, the target

format 'lp', 'opb', or 'smt' may be specified (equal to 'opb' by default). Finally, supplementary comment lines can be specified in the `comments` parameter.

Note: `SMT-LIB2` does not directly support PB constraints. As a result, native cardinality constraints of CNF+ cannot be translated to SMT-LIB2 unless an explicit cardinality encoding is applied. You may want to use Z3's API instead (see its PB interface).

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **format** (*str*) – alien file format to use
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import CNFPlus
>>> cnf = CNFPlus()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.lp', 'w') as fp:
...     cnf.to_alien(fp, format='lp') # writing to the file pointer
```

to_fp(*file_pointer*, *comments=None*)

The method can be used to save a CNF+ formula into a file pointer. The file pointer is expected as an argument. Additionally, supplementary comment lines can be specified in the `comments` parameter.

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import CNFPlus
>>> cnf = CNFPlus()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.cnf+', 'w') as fp:
...     cnf.to_fp(fp) # writing to the file pointer
```

weighted()

This method creates a weighted copy of the internal formula. As a result, an object of class `WCNFPlus` is returned. Every clause of the CNFPlus formula is *soft* in the new WCNFPlus formula and its weight is equal to 1. The set of hard clauses of the new formula is empty. The set of cardinality constraints remains unchanged.

Returns an object of class `WCNFPlus`.

Example:

```
>>> from pysat.formula import CNFPlus
>>> cnf = CNFPlus()
>>> cnf.append([-1, 2])
```

(continues on next page)

(continued from previous page)

```

>>> cnf.append([3, 4])
>>> cnf.append([[1, 2], 1], is_atmost=True)
>>>
>>> wcnf = cnf.weighted()
>>> print(wcnf.hard)
[]
>>> print(wcnf.soft)
[[-1, 2], [3, 4]]
>>> print(wcnf.wght)
[1, 1]
>>> print(wcnf.atms)
[[[1, 2], 1]]

```

class pysat.formula.IDPool(*start_from=1, occupied=[]*)

A simple manager of variable IDs. It can be used as a pool of integers assigning an ID to any object. Identifiers are to start from 1 by default. The list of occupied intervals is empty by default. If necessary the top variable ID can be accessed directly using the `top` variable.

Parameters

- **start_from** (*int*) – the smallest ID to assign.
- **occupied** (*list(list(int))*) – a list of occupied intervals.

id(*obj=None*)

The method is to be used to assign an integer variable ID for a given new object. If the object already has an ID, no new ID is created and the old one is returned instead.

An object can be anything. In some cases it is convenient to use string variable names. Note that if the object is not provided, the method will return a new id unassigned to any object.

Parameters **obj** – an object to assign an ID to.

Return type *int*.

Example:

```

>>> from pysat.formula import IDPool
>>> vpool = IDPool(occupied=[[12, 18], [3, 10]])
>>>
>>> # creating 5 unique variables for the following strings
>>> for i in range(5):
...     print(vpool.id('v{0}'.format(i + 1)))
1
2
11
19
20

```

In some cases, it makes sense to create an external function for accessing IDPool, e.g.:

```

>>> # continuing the previous example
>>> var = lambda i: vpool.id('var{0}'.format(i))
>>> var(5)
20

```

(continues on next page)

(continued from previous page)

```
>>> var('hello_world!')
21
```

obj(vid)

The method can be used to map back a given variable identifier to the original object labeled by the identifier.

Parameters **vid** (*int*) – variable identifier.

Returns an object corresponding to the given identifier.

Example:

```
>>> vpool.obj(21)
'hello_world!'
```

occupy(start, stop)

Mark a given interval as occupied so that the manager could skip the values from **start** to **stop** (**inclusive**).

Parameters

- **start** (*int*) – beginning of the interval.
- **stop** (*int*) – end of the interval.

restart(start_from=1, occupied=[])

Restart the manager from scratch. The arguments replicate those of the constructor of [IDPool](#).

class pysat.formula.WCNF(from_file=None, from_fp=None, from_string=None, comment_lead=['c'])

Class for manipulating partial (weighted) CNF formulas. It can be used for creating formulas, reading them from a file, or writing them to a file. The **comment_lead** parameter can be helpful when one needs to parse specific comment lines starting not with character **c** but with another character or a string.

Parameters

- **from_file** (*str*) – a DIMACS CNF filename to read from
- **from_fp** (*file_pointer*) – a file pointer to read from
- **from_string** (*str*) – a string storing a CNF formula
- **comment_lead** (*list(str)*) – a list of characters leading comment lines

append(clause, weight=None)

Add one more clause to WCNF formula. This method additionally updates the number of variables, i.e. variable **self.nv**, used in the formula.

The clause can be hard or soft depending on the **weight** argument. If no weight is set, the clause is considered to be hard.

Parameters

- **clause** (*list(int)*) – a new clause to add.
- **weight** (*integer or None*) – integer weight of the clause.

```
>>> from pysat.formula import WCNF
>>> cnf = WCNF()
>>> cnf.append([-1, 2])
>>> cnf.append([1], weight=10)
>>> cnf.append([-2], weight=20)
```

(continues on next page)

(continued from previous page)

```
>>> print(cnf.hard)
[[-1, 2]]
>>> print(cnf.soft)
[[1], [-2]]
>>> print(cnf.wght)
[10, 20]
```

copy()

This method can be used for creating a copy of a WCNF object. It creates another object of the *WCNF* class and makes use of the *deepcopy* functionality to copy both hard and soft clauses.

Returns an object of class *WCNF*.

Example:

```
>>> cnf1 = WCNF()
>>> cnf1.append([-1, 2])
>>> cnf1.append([1], weight=10)
>>>
>>> cnf2 = cnf1.copy()
>>> print(cnf2.hard)
[[-1, 2]]
>>> print(cnf2.soft)
[[1]]
>>> print(cnf2.wght)
[10]
>>> print(cnf2.nv)
2
```

extend(*clauses*, *weights=None*)

Add several clauses to WCNF formula. The clauses should be given in the form of list. For every clause in the list, method *append()* is invoked.

The clauses can be hard or soft depending on the *weights* argument. If no weights are set, the clauses are considered to be hard.

Parameters

- **clauses** (*list(list(int))*) – a list of new clauses to add.
- **weights** (*list(int)*) – a list of integer weights.

Example:

```
>>> from pysat.formula import WCNF
>>> cnf = WCNF()
>>> cnf.extend([[ -3, 4], [5, 6]])
>>> cnf.extend([[3], [-4], [-5], [-6]], weights=[1, 5, 3, 4])
>>> print(cnf.hard)
[[-3, 4], [5, 6]]
>>> print(cnf.soft)
[[3], [-4], [-5], [-6]]
>>> print(cnf.wght)
[1, 5, 3, 4]
```

from_file(*fname*, *comment_lead*=['c'], *compressed_with*='use_ext')

Read a WCNF formula from a file in the DIMACS format. A file name is expected as an argument. A default argument is *comment_lead* for parsing comment lines. A given file can be compressed by either gzip, bzip2, or lzma.

Parameters

- **fname** (*str*) – name of a file to parse.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines
- **compressed_with** (*str*) – file compression algorithm

Note that the *compressed_with* parameter can be None (i.e. the file is uncompressed), 'gzip', 'bzip2', 'lzma', or 'use_ext'. The latter value indicates that compression type should be automatically determined based on the file extension. Using 'lzma' in Python 2 requires the `backports.lzma` package to be additionally installed.

Usage example:

```
>>> from pysat.formula import WCNF
>>> cnf1 = WCNF()
>>> cnf1.from_file('some-file.wcnf.bz2', compressed_with='bzip2')
>>>
>>> cnf2 = WCNF(from_file='another-file.wcnf')
```

from_fp(*file_pointer*, *comment_lead*=['c'])

Read a WCNF formula from a file pointer. A file pointer should be specified as an argument. The only default argument is *comment_lead*, which can be used for parsing specific comment lines.

Parameters

- **file_pointer** (*file pointer*) – a file pointer to read the formula from.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines

Usage example:

```
>>> with open('some-file.cnf', 'r') as fp:
...     cnf1 = WCNF()
...     cnf1.from_fp(fp)
>>>
>>> with open('another-file.cnf', 'r') as fp:
...     cnf2 = WCNF(from_fp=fp)
```

from_string(*string*, *comment_lead*=['c'])

Read a WCNF formula from a string. The string should be specified as an argument and should be in the DIMACS CNF format. The only default argument is *comment_lead*, which can be used for parsing specific comment lines.

Parameters

- **string** (*str*) – a string containing the formula in DIMACS.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines

Example:

```
>>> from pysat.formula import WCNF
>>> cnf1 = WCNF()
>>> cnf1.from_string('p wcnf 2 2 2\n 2 -1 2 0\n1 1 -2 0')
>>> print(cnf1.hard)
[[-1, 2]]
>>> print(cnf1.soft)
[[1, 2]]
>>>
>>> cnf2 = WCNF(from_string='p wcnf 3 3 2\n2 -1 2 0\n2 -2 3 0\n1 -3 0\n')
>>> print(cnf2.hard)
[[-1, 2], [-2, 3]]
>>> print(cnf2.soft)
[[-3]]
>>> print(cnf2.nv)
3
```

normalize_negatives(negatives)

Iterate over all soft clauses with negative weights and add their negation either as a hard clause or a soft one.

Parameters **negatives** (*list(list(int))*) – soft clauses with their negative weights.

to_alien(file_pointer, format='opb', comments=None)

The method can be used to dump a WCNF formula into a file pointer in an alien file format, which at this point can either be LP, OPB, or SMT. The file pointer is expected as an argument. Additionally, the target format 'lp', 'opb', or 'smt' may be specified (equal to 'opb' by default). Finally, supplementary comment lines can be specified in the `comments` parameter.

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **format** (*str*) – alien file format to use
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import WCNF
>>> cnf = WCNF()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.lp', 'w') as fp:
...     cnf.to_alien(fp, format='lp') # writing to the file pointer
```

to_file(fname, comments=None, compress_with='use_ext')

The method is for saving a WCNF formula into a file in the DIMACS CNF format. A file name is expected as an argument. Additionally, supplementary comment lines can be specified in the `comments` parameter. Also, a file can be compressed using either gzip, bzip2, or lzma (xz).

Parameters

- **fname** (*str*) – a file name where to store the formula.
- **comments** (*list(str)*) – additional comments to put in the file.
- **compress_with** (*str*) – file compression algorithm

Note that the `compress_with` parameter can be `None` (i.e. the file is uncompressed), `'gzip'`, `'bzip2'`, `'lzma'`, or `'use_ext'`. The latter value indicates that compression type should be automatically determined based on the file extension. Using `'lzma'` in Python 2 requires the `backports.lzma` package to be additionally installed.

Example:

```
>>> from pysat.formula import WCNF
>>> wcnf = WCNF()
...
>>> # the formula is filled with a bunch of clauses
>>> wcnf.to_file('some-file-name.wcnf') # writing to a file
```

to_fp(*file_pointer*, *comments=None*)

The method can be used to save a WCNF formula into a file pointer. The file pointer is expected as an argument. Additionally, supplementary comment lines can be specified in the `comments` parameter.

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import WCNF
>>> wcnf = WCNF()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.wcnf', 'w') as fp:
...     wcnf.to_fp(fp) # writing to the file pointer
```

unweighted()

This method creates a *plain* (unweighted) copy of the internal formula. As a result, an object of class `CNF` is returned. Every clause (both hard or soft) of the WCNF formula is copied to the `clauses` variable of the resulting plain formula, i.e. all weights are discarded.

Returns an object of class `CNF`.

Example:

```
>>> from pysat.formula import WCNF
>>> wcnf = WCNF()
>>> wcnf.extend([-3, 4], [5, 6])
>>> wcnf.extend([3], [-4], [-5], [-6], weights=[1, 5, 3, 4])
>>>
>>> cnf = wcnf.unweighted()
>>> print(cnf.clauses)
[[-3, 4], [5, 6], [3], [-4], [-5], [-6]]
```

class `pysat.formula.WCNFPlus`(*from_file=None*, *from_fp=None*, *from_string=None*, *comment_lead=['c']*)

WCNF formulas augmented with *native* cardinality constraints.

This class inherits most of the functionality of the `WCNF` class. The only difference between the two is that `WCNFPlus` supports *native* cardinality constraints of `MiniCard`.

The parser of input DIMACS files of `WCNFPlus` assumes the syntax of `AtMostK` and `AtLeastK` constraints following the one defined for `CNFPlus` in the [description](#) of `MiniCard`:

```
c Example: Two (hard) cardinality constraints followed by a soft clause
p wcnf+ 7 3 10
10 1 -2 3 5 -7 <= 3
10 4 5 6 -7 >= 2
5 3 5 7 0
```

Note that every cardinality constraint is assumed to be *hard*, i.e. soft cardinality constraints are currently *not supported*.

Each `AtLeastK` constraint is translated into an `AtMostK` constraint in the standard way: $\sum_{i=1}^n x_i \geq k \leftrightarrow \sum_{i=1}^n \neg x_i \leq (n - k)$. Internally, `AtMostK` constraints are stored in variable `atms`, each being a pair `(lits, k)`, where `lits` is a list of literals in the sum and `k` is the upper bound.

Example:

```
>>> from pysat.formula import WCNFPlus
>>> cnf = WCNFPlus(from_string='p wcnf+ 7 3 10\n10 1 -2 3 5 -7 <= 3\n10 4 5 6 -7 >= 2\n5 3 5 7 0\n')
>>> print(cnf.soft)
[[3, 5, 7]]
>>> print(cnf.wght)
[5]
>>> print(cnf.hard)
[]
>>> print(cnf.atms)
[[[1, -2, 3, 5, -7], 3], [[-4, -5, -6, 7], 2]]
>>> print(cnf.nv)
7
```

For details on the functionality, see [WCNF](#).

append(*clause*, *weight=None*, *is_atmost=False*)

Add a single clause or a single `AtMostK` constraint to `WCNF+` formula. This method additionally updates the number of variables, i.e. variable `self.nv`, used in the formula.

If the clause is an `AtMostK` constraint, this should be set with the use of the additional default argument `is_atmost`, which is set to `False` by default.

If `is_atmost` is set to `False`, the clause can be either hard or soft depending on the `weight` argument. If no weight is specified, the clause is considered hard. Otherwise, the clause is soft.

Parameters

- **clause** (*list(int)*) – a new clause to add.
- **weight** (*integer or None*) – an integer weight of the clause.
- **is_atmost** (*bool*) – if `True`, the clause is `AtMostK`.

```
>>> from pysat.formula import WCNFPlus
>>> cnf = WCNFPlus()
>>> cnf.append([-3, 4])
>>> cnf.append([1, 2, 3], 1, is_atmost=True)
>>> cnf.append([-1, -2], weight=35)
>>> print(cnf.hard)
[[-3, 4]]
>>> print(cnf.atms)
```

(continues on next page)

(continued from previous page)

```
[[1, 2, 3], 1]
>>> print(cnf.soft)
[[-1, -2]]
>>> print(cnf.wght)
[35]
```

copy()

This method can be used for creating a copy of a `WCNFPlus` object. It creates another object of the `WCNFPlus` class, call the copy function of `WCNF` class and makes use of the *deepcopy* functionality to copy the atmost constraints.

Returns an object of class `WCNFPlus`.

Example:

```
>>> cnf1 = WCNFPlus()
>>> cnf1.append([-1, 2])
>>> cnf1.append([1], weight=10)
>>> cnf1.append([1, 2], 1, is_atmost=True)
>>> cnf2 = cnf1.copy()
>>> print(cnf2.hard)
[[-1, 2]]
>>> print(cnf2.soft)
[[1]]
>>> print(cnf2.wght)
[10]
>>> print(cnf2.nv)
2
>> print(cnf2.atms)
[[[1, 2], 1]]
```

from_fp(file_pointer, comment_lead=['c'])

Read a `WCNF+` formula from a file pointer. A file pointer should be specified as an argument. The only default argument is `comment_lead`, which can be used for parsing specific comment lines.

Parameters

- **file_pointer** (*file pointer*) – a file pointer to read the formula from.
- **comment_lead** (*list(str)*) – a list of characters leading comment lines

Usage example:

```
>>> with open('some-file.wcnf+', 'r') as fp:
...     cnf1 = WCNFPlus()
...     cnf1.from_fp(fp)
>>>
>>> with open('another-file.wcnf+', 'r') as fp:
...     cnf2 = WCNFPlus(from_fp=fp)
```

to_alien(file_pointer, format='opb', comments=None)

The method can be used to dump a `WCNF+` formula into a file pointer in an alien file format, which at this point can either be LP, OPB, or SMT. The file pointer is expected as an argument. Additionally, the target format 'lp', 'opb', or 'smt' may be specified (equal to 'opb' by default). Finally, supplementary comment lines can be specified in the `comments` parameter.

Note: `SMT-LIB2` does not directly support PB constraints. As a result, native cardinality constraints of CNF+ cannot be translated to SMT-LIB2 unless an explicit cardinality encoding is applied. You may want to use Z3's API instead (see its PB interface).

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **format** (*str*) – alien file format to use
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import WCNFPlus
>>> cnf = WCNFPlus()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.lp', 'w') as fp:
...     cnf.to_alien(fp, format='lp') # writing to the file pointer
```

to_fp(*file_pointer*, *comments=None*)

The method can be used to save a WCNF+ formula into a file pointer. The file pointer is expected as an argument. Additionally, supplementary comment lines can be specified in the `comments` parameter.

Parameters

- **file_pointer** (*file pointer*) – a file pointer where to store the formula.
- **comments** (*list(str)*) – additional comments to put in the file.

Example:

```
>>> from pysat.formula import WCNFPlus
>>> cnf = WCNFPlus()
...
>>> # the formula is filled with a bunch of clauses
>>> with open('some-file.wcnf+', 'w') as fp:
...     cnf.to_fp(fp) # writing to the file pointer
```

unweighted()

This method creates a *plain* (unweighted) copy of the internal formula. As a result, an object of class `CNFPlus` is returned. Every clause (both hard or soft) of the original WCNFPlus formula is copied to the `clauses` variable of the resulting plain formula, i.e. all weights are discarded.

Note that the cardinality constraints of the original (weighted) formula remain unchanged in the new (plain) formula.

Returns an object of class `CNFPlus`.

Example:

```
>>> from pysat.formula import WCNF
>>> wcnf = WCNFPlus()
>>> wcnf.extend([[3, 4], [5, 6]])
>>> wcnf.extend([[3], [-4], [-5], [-6]], weights=[1, 5, 3, 4])
```

(continues on next page)

(continued from previous page)

```

>>> wcnf.append([[1, 2, 3], 1], is_atmost=True)
>>>
>>> cnf = wcnf.unweighted()
>>> print(cnf.clauses)
[[-3, 4], [5, 6], [3], [-4], [-5], [-6]]
>>> print(cnf.atmosts)
[[[1, 2, 3], 1]]

```

1.1.3 Pseudo-Boolean encodings (pysat.pb)

List of classes

<i>EncType</i>	This class represents a C-like enum type for choosing the pseudo-Boolean encoding to use.
<i>PBEnc</i>	Abstract class responsible for the creation of pseudo-Boolean constraints encoded to a CNF formula.

Module description

Note: Functionality of this module is available only if the *PyPBLib* package is installed, e.g. from PyPI:

```
$ pip install pypblib
```

This module provides access to the basic functionality of the *PyPBLib* library developed by the *Logic Optimization Group* of the University of Lleida. *PyPBLib* provides a user with an extensive Python API to the well-known *PBLib* library¹. Note the *PyPBLib* has a number of *additional features* that cannot be accessed through *PySAT* *at this point*. (One concrete example is a range of cardinality encodings, which clash with the internal *pysat.card* module.) If a user needs some functionality of *PyPBLib* missing in this module, he/she may apply *PyPBLib* as a standalone library, when working with *PySAT*.

A *pseudo-Boolean constraint* is a constraint of the form: $(\sum_{i=1}^n a_i \cdot x_i) \circ k$, where $a_i \in \mathbb{N}$, $x_i \in \{y_i, \neg y_i\}$, $y_i \in \mathbb{B}$, and $\circ \in \{\leq, =, \geq\}$. Pseudo-Boolean constraints arise in a number of important practical applications. Thus, several *encodings* of pseudo-Boolean constraints into CNF formulas are known². The list of pseudo-Boolean encodings supported by this module include BDD³⁴, sequential weight counters⁵, sorting networks³, adder networks³, and binary merge⁶. Access to all cardinality encodings can be made through the main class of this module, which is *PBEnc*.

¹ Tobias Philipp, Peter Steinke. *PBLib - A Library for Encoding Pseudo-Boolean Constraints into CNF*. SAT 2015. pp. 9-16

² Olivier Roussel, Vasco M. Manquinho. *Pseudo-Boolean and Cardinality Constraints*. Handbook of Satisfiability. 2009. pp. 695-733

³ Niklas Eén, Niklas Sörensson. *Translating Pseudo-Boolean Constraints into SAT*. JSAT. vol. 2(1-4). 2006. pp. 1-26

⁴ Ignasi Abío, Robert Nieuwenhuis, Albert Oliveras, Enric Rodríguez-Carbonell. *BDDs for Pseudo-Boolean Constraints - Revisited*. SAT. 2011. pp. 61-75

⁵ Steffen Hölldobler, Norbert Manthey, Peter Steinke. *A Compact Encoding of Pseudo-Boolean Constraints into SAT*. KI. 2012. pp. 107-118

⁶ Norbert Manthey, Tobias Philipp, Peter Steinke. *A More Compact Translation of Pseudo-Boolean Constraints into CNF Such That Generalized Arc Consistency Is Maintained*. KI. 2014. pp. 123-134

Module details

class pysat.pb.EncType

This class represents a C-like enum type for choosing the pseudo-Boolean encoding to use. The values denoting the encodings are:

```
best      = 0
bdd       = 1
seqcounter = 2
sortnetwrk = 3
adder     = 4
binmerge  = 5
```

The desired encoding can be selected either directly by its integer identifier, e.g. 2, or by its alphabetical name, e.g. `EncType.seqcounter`.

All the encodings are produced and returned as a list of clauses in the `pysat.formula.CNF` format.

Note that the encoding type can be set to `best`, in which case the encoder selects one of the other encodings from the list (in most cases, this invokes the `bdd` encoder).

exception pysat.pb.NoSuchEncodingError

This exception is raised when creating an unknown LEQ, GEQ, or Equals constraint encoding.

with_traceback()

Exception.with_traceback(tb) – set self.__traceback__ to tb and return self.

class pysat.pb.PBEnc

Abstract class responsible for the creation of pseudo-Boolean constraints encoded to a CNF formula. The class has three main *class methods* for creating LEQ, GEQ, and Equals constraints. Given (1) either a list of weighted literals or a list of unweighted literals followed by a list of weights, (2) an integer bound and an encoding type, each of these methods returns an object of class `pysat.formula.CNF` representing the resulting CNF formula.

Since the class is abstract, there is no need to create an object of it. Instead, the methods should be called directly as class methods, e.g. `PBEnc.atmost(wlits, bound)` or `PBEnc.equals(lits, weights, bound)`. An example usage is the following:

```
>>> from pysat.pb import *
>>> cnf = PBEnc.atmost(lits=[1, 2, 3], weights=[1, 2, 3], bound=3)
>>> print(cnf.clauses)
[[4], [-1, -5], [-2, -5], [5, -3, -6], [6]]
>>> cnf = PBEnc.equals(lits=[1, 2, 3], weights=[1, 2, 3], bound=3, encoding=EncType.
↳ bdd)
>>> print(cnf.clauses)
[[4], [-5, -2], [-5, 2, -1], [-5, -1], [-6, 1], [-7, -2, 6], [-7, 2], [-7, 6], [-8,
↳ -3, 5], [-8, 3, 7], [-8, 5, 7], [8]]
```

classmethod atleast(lits, weights=None, bound=1, top_id=None, vpool=None, encoding=0)

A synonym for `PBEnc.geq()`.

classmethod atmost(lits, weights=None, bound=1, top_id=None, vpool=None, encoding=0)

A synonym for `PBEnc.leq()`.

classmethod equals(lits, weights=None, bound=1, top_id=None, vpool=None, encoding=0)

This method can be used for creating a CNF encoding of a weighted EqualsK constraint, i.e. of $\sum_{i=1}^n a_i \cdot x_i = k$. The method shares the arguments and the return type with method `PBEnc.leq()`. Please, see it for details.

classmethod `geq(lits, weights=None, bound=1, top_id=None, vpool=None, encoding=0)`

This method can be used for creating a CNF encoding of a GEQ (weighted AtLeastK) constraint, i.e. of $\sum_{i=1}^n a_i \cdot x_i \geq k$. The method shares the arguments and the return type with method `PBEnc.leq()`. Please, see it for details.

classmethod `leq(lits, weights=None, bound=1, top_id=None, vpool=None, encoding=0)`

This method can be used for creating a CNF encoding of a LEQ (weighted AtMostK) constraint, i.e. of $\sum_{i=1}^n a_i \cdot x_i \leq k$. The resulting set of clauses is returned as an object of class `formula.CNF`.

The input list of literals can contain either integers or pairs (l, w) , where l is an integer literal and w is an integer weight. The latter can be done only if no `weights` are specified separately. The type of encoding to use can be specified using the `encoding` parameter. By default, it is set to `EncType.best`, i.e. it is up to the PBLib encoder to choose the encoding type.

Parameters

- **`lits`** (*iterable(int)*) – a list of literals in the sum.
- **`weights`** (*iterable(int)*) – a list of weights
- **`bound`** (*int*) – the value of bound k .
- **`top_id`** (*integer or None*) – top variable identifier used so far.
- **`vpool`** (*IDPool*) – variable pool for counting the number of variables.
- **`encoding`** (*integer*) – identifier of the encoding to use.

Return type `pysat.formula.CNF`

1.1.4 SAT solvers' API (`pysat.solvers`)

List of classes

<i>SolverNames</i>	This class serves to determine the solver requested by a user given a string name.
<i>Solver</i>	Main class for creating and manipulating a SAT solver.
Cadical	CaDiCaL SAT solver.
Gluecard3	Gluecard 3 SAT solver.
Gluecard4	Gluecard 4 SAT solver.
Glucose3	Glucose 3 SAT solver.
Glucose4	Glucose 4.1 SAT solver.
Lingeling	Lingeling SAT solver.
MapleChrono	MapleLCMDistChronoBT SAT solver.
MapleCM	MapleCM SAT solver.
Maplesat	MapleCOMSPS_LRB SAT solver.
Mergesat3	MergeSat 3 SAT solver.
Minicard	Minicard SAT solver.
Minisat22	MiniSat 2.2 SAT solver.
MinisatGH	MiniSat SAT solver (version from github).

Module description

This module provides *incremental* access to a few modern SAT solvers. The solvers supported by PySAT are:

- CaDiCaL ([rel-1.0.3](#))
- Glucose (3.0)
- Glucose (4.1)
- Lingeling ([bbc-9230380-160707](#))
- MapleLCMDistChronoBT ([SAT competition 2018 version](#))
- MapleCM ([SAT competition 2018 version](#))
- Maplesat ([MapleCOMSPS_LRB](#))
- Mergesat (3.0)
- Minicard (1.2)
- Minisat ([2.2 release](#))
- Minisat ([GitHub version](#))

Additionally, PySAT includes the versions of Glucose3 and Glucose4 that support native cardinality constraints, ported from Minicard:

- Gluecard3
- Gluecard4

All solvers can be accessed through a unified MiniSat-like¹ incremental² interface described below.

The module provides direct access to all supported solvers using the corresponding classes `Cadical`, `Gluecard3`, `Gluecard4`, `Glucose3`, `Glucose4`, `Lingeling`, `MapleChrono`, `MapleCM`, `Maplesat`, `Mergesat3`, `Minicard`, `Minisat22`, and `MinisatGH`. However, the solvers can also be accessed through the common base class `Solver` using the solver `name` argument. For example, both of the following pieces of code create a copy of the `Glucose3` solver:

```
>>> from pysat.solvers import Glucose3, Solver
>>>
>>> g = Glucose3()
>>> g.delete()
>>>
>>> s = Solver(name='g3')
>>> s.delete()
```

The `pysat.solvers` module is designed to create and manipulate SAT solvers as *oracles*, i.e. it does not give access to solvers' internal parameters such as variable polarities or activities. PySAT provides a user with the following basic SAT solving functionality:

- creating and deleting solver objects
- adding individual clauses and formulas to solver objects
- making SAT calls with or without assumptions
- propagating a given set of assumption literals
- setting preferred polarities for a (sub)set of variables
- extracting a model of a satisfiable input formula

¹ Niklas Eén, Niklas Sörensson. *An Extensible SAT-solver*. SAT 2003, pp. 502-518

² Niklas Eén, Niklas Sörensson. *Temporal induction by incremental SAT solving*. Electr. Notes Theor. Comput. Sci. 89(4). 2003. pp. 543-560

- enumerating models of an input formula
- extracting an unsatisfiable core of an unsatisfiable formula
- extracting a **DRUP proof** logged by the solver

PySAT supports both non-incremental and incremental SAT solving. Incrementality can be achieved with the use of the MiniSat-like *assumption-based* interface^{Page 32, 2}. It can be helpful if multiple calls to a SAT solver are needed for the same formula using different sets of “assumptions”, e.g. when doing consecutive SAT calls for formula $\mathcal{F} \wedge (a_{i_1} \wedge \dots \wedge a_{i_1+j_1})$ and $\mathcal{F} \wedge (a_{i_2} \wedge \dots \wedge a_{i_2+j_2})$, where every a_{i_k} is an assumption literal.

There are several advantages of using assumptions: (1) it enables one to *keep and reuse* the clauses learnt during previous SAT calls at a later stage and (2) assumptions can be easily used to extract an *unsatisfiable core* of the formula. A drawback of assumption-based SAT solving is that the clauses learnt are longer (they typically contain many assumption literals), which makes the SAT calls harder.

In PySAT, assumptions should be provided as a list of literals given to the `solve()` method:

```
>>> from pysat.solvers import Solver
>>> s = Solver()
>>>
... # assume that solver s is fed with a formula
>>>
>>> s.solve() # a simple SAT call
True
>>>
>>> s.solve(assumptions=[1, -2, 3]) # a SAT call with assumption literals
False
>>> s.get_core() # extracting an unsatisfiable core
[3, 1]
```

In order to shorten the description of the module, the classes providing direct access to the individual solvers, i.e. classes `Cadical`, `Gluecard3`, `Gluecard4`, `Glucose3`, `Glucose4`, `Lingeling`, `MapleChrono`, `MapleCM`, `Maplesat`, `Mergesat3`, `Minicard`, `Minisat22`, and `MinisatGH`, are **omitted**. They replicate the interface of the base class `Solver` and, thus, can be used the same exact way.

Module details

exception `pysat.solvers.NoSuchSolverError`

This exception is raised when creating a new SAT solver whose name does not match any name in `SolverNames`. The list of *known* solvers includes the names ‘`cadical`’, ‘`gluecard3`’, ‘`gluecard4`’, ‘`glucose3`’, ‘`glucose4`’, ‘`lingeling`’, ‘`maplechrono`’, ‘`maplecm`’, ‘`maplesat`’, ‘`mergesat3`’, ‘`minicard`’, ‘`minisat22`’, and ‘`minisatgh`’.

`with_traceback()`

`Exception.with_traceback(tb)` – set `self.__traceback__` to `tb` and return `self`.

class `pysat.solvers.Solver(name='m22', bootstrap_with=None, use_timer=False, **kwargs)`

Main class for creating and manipulating a SAT solver. Any available SAT solver can be accessed as an object of this class and so `Solver` can be seen as a wrapper for all supported solvers.

The constructor of `Solver` has only one mandatory argument `name`, while all the others are default. This means that explicit solver constructors, e.g. `Glucose3` or `MinisatGH` etc., have only default arguments.

Parameters

- **name** (*str*) – solver’s name (see `SolverNames`).
- **bootstrap_with** (*iterable(iterable(int))*) – a list of clauses for solver initialization.

- **use_timer** (*bool*) – whether or not to measure SAT solving time.

The `bootstrap_with` argument is useful when there is an input CNF formula to feed the solver with. The argument expects a list of clauses, each clause being a list of literals, i.e. a list of integers.

If set to `True`, the `use_timer` parameter will force the solver to accumulate the time spent by all SAT calls made with this solver but also to keep time of the last SAT call.

Once created and used, a solver must be deleted with the `delete()` method. Alternatively, if created using the `with` statement, deletion is done automatically when the end of the `with` block is reached.

Given the above, a couple of examples of solver creation are the following:

```
>>> from pysat.solvers import Solver, Minisat22
>>>
>>> s = Solver(name='g4')
>>> s.add_clause([-1, 2])
>>> s.add_clause([-1, -2])
>>> s.solve()
True
>>> print(s.get_model())
[-1, -2]
>>> s.delete()
>>>
>>> with Minisat22(bootstrap_with=[[-1, 2], [-1, -2]]) as m:
...     m.solve()
True
...     print(m.get_model())
[-1, -2]
```

Note that while all explicit solver classes necessarily have default arguments `bootstrap_with` and `use_timer`, solvers `Cadical`, `Lingeling`, `Gluecard3`, `Gluecard4`, `Glucose3`, `Glucose4`, `MapleChrono`, `MapleCM`, and `Maplesat` can have additional default arguments. One such argument supported by is `DRUP proof` logging. This can be enabled by setting the `with_proof` argument to `True` (`False` by default):

```
>>> from pysat.solvers import Lingeling
>>> from pysat.examples.genhard import PHP
>>>
>>> cnf = PHP(nof_holes=2) # pigeonhole principle for 3 pigeons
>>>
>>> with Lingeling(bootstrap_with=cnf.clauses, with_proof=True) as l:
...     l.solve()
False
...     l.get_proof()
['-5 0', '6 0', '-2 0', '-4 0', '1 0', '3 0', '0']
```

Additionally, Glucose-based solvers, namely `Glucose3`, `Glucose4`, `Gluecard3`, and `Gluecard4` have one more default argument `incr` (`False` by default), which enables incrementality features introduced in `Glucose3`³. To summarize, the additional arguments of Glucose are:

Parameters

- **incr** (*bool*) – enable the incrementality features of `Glucose3`³.
- **with_proof** (*bool*) – enable proof logging in the `DRUP` format.

³ Gilles Audemard, Jean-Marie Lagniez, Laurent Simon. *Improving Glucose for Incremental SAT Solving with Assumptions: Application to MUS Extraction*. SAT 2013. pp. 309-317

Finally, most MiniSat-based solvers can be exploited in the “warm-start” mode in the case of *satisfiable* formulas. This may come in handy in various model enumeration settings. Note that warm-start mode is disabled in the case of limited solving with “*unknown*” outcomes. Warm-start mode can be set with the use of the *warm_start* parameter:

Parameters *warm_start* (*bool*) – use the solver in the “warm-start” mode

accum_stats()

Get accumulated low-level stats from the solver. Currently, the statistics includes the number of restarts, conflicts, decisions, and propagations.

Return type dict.

Example:

```
>>> from pysat.examples.genhard import PHP
>>> cnf = PHP(5)
>>> from pysat.solvers import Solver
>>> with Solver(bootstrap_with=cnf) as s:
...     print(s.solve())
...     print(s.accum_stats())
False
{'restarts': 2, 'conflicts': 201, 'decisions': 254, 'propagations': 2321}
```

add_atmost(*lits*, *k*, *no_return=True*)

This method is responsible for adding a new *native* AtMostK (see *pysat.card*) constraint.

Note that most of the solvers do not support native AtMostK constraints.

An AtMostK constraint is $\sum_{i=1}^n x_i \leq k$. A native AtMostK constraint should be given as a pair *lits* and *k*, where *lits* is a list of literals in the sum.

Parameters

- **lits** (*iterable(int)*) – a list of literals.
- **k** (*int*) – upper bound on the number of satisfied literals
- **no_return** (*bool*) – check solver’s internal formula and return the result, if set to False.

Return type bool if *no_return* is set to False.

A usage example is the following:

```
>>> s = Solver(name='mc', bootstrap_with=[[1], [2], [3]])
>>> s.add_atmost(lits=[1, 2, 3], k=2, no_return=False)
False
>>> # the AtMostK constraint is in conflict with initial unit clauses
```

add_clause(*clause*, *no_return=True*)

This method is used to add a single clause to the solver. An optional argument *no_return* controls whether or not to check the formula’s satisfiability after adding the new clause.

Parameters

- **clause** (*iterable(int)*) – an iterable over literals.
- **no_return** (*bool*) – check solver’s internal formula and return the result, if set to False.

Return type bool if *no_return* is set to False.

Note that a clause can be either a list of integers or another iterable type over integers, e.g. tuple or set among others.

A usage example is the following:

```
>>> s = Solver(bootstrap_with=[[-1, 2], [-1, -2]])
>>> s.add_clause([1], no_return=False)
False
```

append_formula(*formula*, *no_return=True*)

This method can be used to add a given list of clauses into the solver.

Parameters

- **formula** (*iterable(iterable(int))*) – a list of clauses.
- **no_return** (*bool*) – check solver’s internal formula and return the result, if set to False.

The *no_return* argument is set to True by default.

Return type bool if *no_return* is set to False.

```
>>> cnf = CNF()
... # assume the formula contains clauses
>>> s = Solver()
>>> s.append_formula(cnf.clauses, no_return=False)
True
```

clear_interrupt()

Clears a previous interrupt. If a limited SAT call was interrupted using the *interrupt()* method, this method **must be called** before calling the SAT solver again.

conf_budget(*budget=-1*)

Set limit (i.e. the upper bound) on the number of conflicts in the next limited SAT call (see *solve_limited()*). The limit value is given as a budget variable and is an integer greater than 0. If the budget is set to 0 or -1, the upper bound on the number of conflicts is disabled.

Parameters **budget** (*int*) – the upper bound on the number of conflicts.

Example:

```
>>> from pysat.solvers import MinisatGH
>>> from pysat.examples.genhard import PHP
>>>
>>> cnf = PHP(nof_holes=20) # PHP20 is too hard for a SAT solver
>>> m = MinisatGH(bootstrap_with=cnf.clauses)
>>>
>>> m.conf_budget(2000) # getting at most 2000 conflicts
>>> print(m.solve_limited()) # making a limited oracle call
None
>>> m.delete()
```

delete()

Solver destructor, which must be called explicitly if the solver is to be removed. This is not needed inside an with block.

enum_models(*assumptions=[]*)

This method can be used to enumerate models of a CNF formula and it performs as a standard Python

iterator. The method can be called without arguments but also with an argument `assumptions`, which represents a list of literals to “assume”.

Warning: Once finished, model enumeration results in the target formula being *unsatisfiable*. This is because the enumeration process *blocks* each previously computed model by adding a new clause until no more models of the formula exist.

Parameters `assumptions` (`iterable(int)`) – a list of assumption literals.

Return type `list(int)`.

Example:

```
>>> with Solver(bootstrap_with=[[-1, 2], [-2, 3]]) as s:
...     for m in s.enum_models():
...         print(m)
[-1, -2, -3]
[-1, -2, 3]
[-1, 2, 3]
[1, 2, 3]
>>>
>>> with Solver(bootstrap_with=[[-1, 2], [-2, 3]]) as s:
...     for m in s.enum_models(assumptions=[1]):
...         print(m)
[1, 2, 3]
```

`get_core()`

This method is to be used for extracting an unsatisfiable core in the form of a subset of a given set of assumption literals, which are responsible for unsatisfiability of the formula. This can be done only if the previous SAT call returned `False` (*UNSAT*). Otherwise, `None` is returned.

Return type `list(int)` or `None`.

Usage example:

```
>>> from pysat.solvers import Minisat22
>>> m = Minisat22()
>>> m.add_clause([-1, 2])
>>> m.add_clause([-2, 3])
>>> m.add_clause([-3, 4])
>>> m.solve(assumptions=[1, 2, 3, -4])
False
>>> print(m.get_core()) # literals 2 and 3 are not in the core
[-4, 1]
>>> m.delete()
```

`get_model()`

The method is to be used for extracting a satisfying assignment for a CNF formula given to the solver. A model is provided if a previous SAT call returned `True`. Otherwise, `None` is reported.

Return type `list(int)` or `None`.

Example:

```
>>> from pysat.solvers import Solver
>>> s = Solver()
>>> s.add_clause([-1, 2])
>>> s.add_clause([-1, -2])
>>> s.add_clause([1, -2])
>>> s.solve()
True
>>> print(s.get_model())
[-1, -2]
>>> s.delete()
```

get_proof()

A DRUP proof can be extracted using this method if the solver was set up to provide a proof. Otherwise, the method returns `None`.

Return type `list(str)` or `None`.

Example:

```
>>> from pysat.solvers import Solver
>>> from pysat.examples.genhard import PHP
>>>
>>> cnf = PHP(nof_holes=3)
>>> with Solver(name='g4', with_proof=True) as g:
...     g.append_formula(cnf.clauses)
...     g.solve()
False
...     print(g.get_proof())
['-8 4 1 0', '-10 0', '-2 0', '-4 0', '-8 0', '-6 0', '0']
```

get_status()

The result of a previous SAT call is stored in an internal variable and can be later obtained using this method.

Return type `Boolean` or `None`.

`None` is returned if a previous SAT call was interrupted.

interrupt()

Interrupt the execution of the current *limited* SAT call (see [solve_limited\(\)](#)). Can be used to enforce time limits using timer objects. The interrupt must be cleared before performing another SAT call (see [clear_interrupt\(\)](#)).

Note that this method can be called if limited SAT calls are made with the option `expect_interrupt` set to `True`.

Behaviour is **undefined** if used to interrupt a *non-limited* SAT call (see [solve\(\)](#)).

Example:

```
>>> from pysat.solvers import MinisatGH
>>> from pysat.examples.genhard import PHP
>>> from threading import Timer
>>>
>>> cnf = PHP(nof_holes=20) # PHP20 is too hard for a SAT solver
>>> m = MinisatGH(bootstrap_with=cnf.clauses)
>>>
```

(continues on next page)

(continued from previous page)

```

>>> def interrupt(s):
>>>     s.interrupt()
>>>
>>> timer = Timer(10, interrupt, [m])
>>> timer.start()
>>>
>>> print(m.solve_limited(expect_interrupt=True))
None
>>> m.delete()

```

new(name='m22', bootstrap_with=None, use_timer=False, **kwargs)

The actual solver constructor invoked from `__init__()`. Chooses the solver to run, based on its name. See [Solver](#) for the parameters description.

Raises [NoSuchSolverError](#) – if there is no solver matching the given name.

nof_clauses()

This method returns the number of clauses currently appearing in the formula given to the solver.

Return type int.

Example:

```

>>> s = Solver(bootstrap_with=[[-1, 2], [-2, 3]])
>>> s.nof_clauses()
2

```

nof_vars()

This method returns the number of variables currently appearing in the formula given to the solver.

Return type int.

Example:

```

>>> s = Solver(bootstrap_with=[[-1, 2], [-2, 3]])
>>> s.nof_vars()
3

```

prop_budget(budget=-1)

Set limit (i.e. the upper bound) on the number of propagations in the next limited SAT call (see [solve_limited\(\)](#)). The limit value is given as a budget variable and is an integer greater than 0. If the budget is set to 0 or -1, the upper bound on the number of conflicts is disabled.

Parameters **budget** (int) – the upper bound on the number of propagations.

Example:

```

>>> from pysat.solvers import MinisatGH
>>> from pysat.examples.genhard import Parity
>>>
>>> cnf = Parity(size=10) # too hard for a SAT solver
>>> m = MinisatGH(bootstrap_with=cnf.clauses)
>>>
>>> m.prop_budget(100000) # doing at most 100000 propagations
>>> print(m.solve_limited()) # making a limited oracle call

```

(continues on next page)

(continued from previous page)

```
None
>>> m.delete()
```

propagate(*assumptions*=[], *phase_saving*=0)

The method takes a list of assumption literals and does unit propagation of each of these literals consecutively. A Boolean status is returned followed by a list of assigned (assumed and also propagated) literals. The status is `True` if no conflict arised during propagation. Otherwise, the status is `False`. Additionally, a user may specify an optional argument `phase_saving` (0 by default) to enable MiniSat-like phase saving.

Note that only MiniSat-like solvers support this functionality (e.g. Cadical and Lingeling do not support it).

Parameters

- **assumptions** (*iterable(int)*) – a list of assumption literals.
- **phase_saving** (*int*) – enable phase saving (can be 0, 1, and 2).

Return type tuple(bool, list(int)).

Usage example:

```
>>> from pysat.solvers import Glucose3
>>> from pysat.card import *
>>>
>>> cnf = CardEnc.atmost(lits=range(1, 6), bound=1, encoding=EncType.pairwise)
>>> g = Glucose3(bootstrap_with=cnf.clauses)
>>>
>>> g.propagate(assumptions=[1])
(True, [1, -2, -3, -4, -5])
>>>
>>> g.add_clause([2])
>>> g.propagate(assumptions=[1])
(False, [])
>>>
>>> g.delete()
```

set_phases(*literals*=[])

The method takes a list of literals as an argument and sets *phases* (or MiniSat-like *polarities*) of the corresponding variables respecting the literals. For example, if a given list of literals is `[1, -513]`, the solver will try to set variable x_1 to true while setting x_{513} to false.

Note that once these preferences are specified, MinisatGH and Lingeling will always respect them when branching on these variables. However, solvers Glucose3, Glucose4, MapleChrono, MapleCM, Maplesat, Minisat22, and Minicard can redefine the preferences in any of the following SAT calls due to the phase saving heuristic.

Also **note** that Cadical does not support this functionality.

Parameters **literals** (*iterable(int)*) – a list of literals.

Usage example:

```
>>> from pysat.solvers import Glucose3
>>>
>>> g = Glucose3(bootstrap_with=[[1, 2]])
>>> # the formula has 3 models: [-1, 2], [1, -2], [1, 2]
```

(continues on next page)

(continued from previous page)

```

>>>
>>> g.set_phases(literals=[1, 2])
>>> g.solve()
True
>>> g.get_model()
[1, 2]
>>>
>>> g.delete()

```

solve(assumptions=[])

This method is used to check satisfiability of a CNF formula given to the solver (see methods [add_clause\(\)](#) and [append_formula\(\)](#)). Unless interrupted with SIGINT, the method returns either True or False.

Incremental SAT calls can be made with the use of assumption literals. (**Note** that the `assumptions` argument is optional and disabled by default.)

Parameters `assumptions` (*iterable(int)*) – a list of assumption literals.

Return type Boolean or None.

Example:

```

>>> from pysat.solvers import Solver
>>> s = Solver(bootstrap_with=[-1, 2], [-2, 3])
>>> s.solve()
True
>>> s.solve(assumptions=[1, -3])
False
>>> s.delete()

```

solve_limited(assumptions=[], expect_interrupt=False)

This method is used to check satisfiability of a CNF formula given to the solver (see methods [add_clause\(\)](#) and [append_formula\(\)](#)), taking into account the upper bounds on the *number of conflicts* (see [conf_budget\(\)](#)) and the *number of propagations* (see [prop_budget\(\)](#)). If the number of conflicts or propagations is set to be larger than 0 then the following SAT call done with [solve_limited\(\)](#) will not exceed these values, i.e. it will be *incomplete*. Otherwise, such a call will be identical to [solve\(\)](#).

As soon as the given upper bound on the number of conflicts or propagations is reached, the SAT call is dropped returning None, i.e. *unknown*. None can also be returned if the call is interrupted by SIGINT. Otherwise, the method returns True or False.

Note that only MiniSat-like solvers support this functionality (e.g. Cadical and Lingeling do not support it).

Incremental SAT calls can be made with the use of assumption literals. (**Note** that the `assumptions` argument is optional and disabled by default.)

Note that since SIGINT handling and [interrupt\(\)](#) are not configured to work *together* at this point, additional input parameter `expect_interrupt` is assumed to be given, indicating what kind of interruption may happen during the execution of [solve_limited\(\)](#): whether a SIGINT signal or internal [interrupt\(\)](#). By default, a SIGINT signal handling is assumed. If `expect_interrupt` is set to True and eventually a SIGINT is received, the behavior is **undefined**.

Parameters

- `assumptions` (*iterable(int)*) – a list of assumption literals.

- **expect_interrupt** (*bool*) – whether *interrupt()* will be called

Return type Boolean or None.

Doing limited SAT calls can be of help if it is known that *complete* SAT calls are too expensive. For instance, it can be useful when minimizing unsatisfiable cores in MaxSAT (see `pysat.examples.RC2.minimize_core()` also shown below).

Also and besides supporting deterministic interruption based on *conf_budget()* and/or *prop_budget()*, limited SAT calls support *deterministic* and *non-deterministic* interruption from inside a Python script. See the *interrupt()* and *clear_interrupt()* methods for details.

Usage example:

```
... # assume that a SAT oracle is set up to contain an unsatisfiable
... # formula, and its core is stored in variable "core"
oracle.conf_budget(1000) # getting at most 1000 conflicts be call

i = 0
while i < len(core):
    to_test = core[:i] + core[(i + 1):]

    # doing a limited call
    if oracle.solve_limited(assumptions=to_test) == False:
        core = to_test
    else: # True or *unknown*
        i += 1
```

supports_atmost()

This method can be called to determine whether the solver supports native AtMostK (see *pysat.card*) constraints.

Return type bool

A usage example is the following:

```
>>> s = Solver(name='mc')
>>> s.supports_atmost()
True
>>> # there is support for AtMostK constraints in this solver
```

time()

Get the time spent when doing the last SAT call. **Note** that the time is measured only if the *use_timer* argument was previously set to True when creating the solver (see *Solver* for details).

Return type float.

Example usage:

```
>>> from pysat.solvers import Solver
>>> from pysat.examples.genhard import PHP
>>>
>>> cnf = PHP(nof_holes=10)
>>> with Solver(bootstrap_with=cnf.clauses, use_timer=True) as s:
...     print(s.solve())
False
...     print('{0:.2f}s'.format(s.time()))
150.16s
```

time_accum()

Get the time spent for doing all SAT calls accumulated. **Note** that the time is measured only if the `use_timer` argument was previously set to `True` when creating the solver (see [Solver](#) for details).

Return type float.

Example usage:

```
>>> from pysat.solvers import Solver
>>> from pysat.examples.genhard import PHP
>>>
>>> cnf = PHP(nof_holes=10)
>>> with Solver(bootstrap_with=cnf.clauses, use_timer=True) as s:
...     print(s.solve(assumptions=[1]))
False
...     print('{0:.2f}s'.format(s.time()))
1.76s
...     print(s.solve(assumptions=[-1]))
False
...     print('{0:.2f}s'.format(s.time()))
113.58s
...     print('{0:.2f}s'.format(s.time_accum()))
115.34s
```

class pysat.solvers.SolverNames

This class serves to determine the solver requested by a user given a string name. This allows for using several possible names for specifying a solver.

```
cadical      = ('cd', 'cdl', 'cadical')
gluecard3    = ('gc3', 'gc30', 'gluecard3', 'gluecard30')
gluecard41   = ('gc4', 'gc41', 'gluecard4', 'gluecard41')
glucose3     = ('g3', 'g30', 'glucose3', 'glucose30')
glucose4     = ('g4', 'g41', 'glucose4', 'glucose41')
lingeling    = ('lgl', 'lingeling')
maplechrono  = ('mcb', 'chrono', 'maplechrono')
maplecm      = ('mcm', 'maplecm')
maplesat     = ('mpl', 'maple', 'maplesat')
mergesat3    = ('mg3', 'mgs3', 'mergesat3', 'mergesat30')
minicard     = ('mc', 'mcard', 'minicard')
minisat22    = ('m22', 'msat22', 'minisat22')
minisatgh    = ('mgh', 'msat-gh', 'minisat-gh')
```

As a result, in order to select Glucose3, a user can specify the solver's name: either 'g3', 'g30', 'glucose3', or 'glucose30'. *Note that the capitalized versions of these names are also allowed.*

1.2 Supplementary examples package

1.2.1 Fu&Malik MaxSAT algorithm (`pysat.examples.fm`)

List of classes

<i>FM</i>	A non-incremental implementation of the FM (Fu&Malik, or WMSU1) algorithm.
-----------	--

Module description

This module implements a variant of the seminal core-guided MaxSAT algorithm originally proposed by¹ and then improved and modified further in²³⁴⁵. Namely, the implementation follows the WMSU1 variant⁵ of the algorithm extended to the case of *weighted partial* formulas.

The implementation can be used as an executable (the list of available command-line options can be shown using `fm.py -h`) in the following way:

```
$ xzcat formula.wcnf.xz
p wcnf 3 6 4
1 1 0
1 2 0
1 3 0
4 -1 -2 0
4 -1 -3 0
4 -2 -3 0

$ fm.py -c cardn -s glucose3 -vv formula.wcnf.xz
c cost: 1; core sz: 2
c cost: 2; core sz: 3
s OPTIMUM FOUND
o 2
v -1 -2 3 0
c oracle time: 0.0001
```

Alternatively, the algorithm can be accessed and invoked through the standard `import` interface of Python, e.g.

```
>>> from pysat.examples.fm import FM
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF(from_file='formula.wcnf.xz')
>>>
>>> fm = FM(wcnf, verbose=0)
>>> fm.compute() # set of hard clauses should be satisfiable
True
>>> print(fm.cost) # cost of MaxSAT solution should be 2
```

(continues on next page)

¹ Zhaohui Fu, Sharad Malik. *On Solving the Partial MAX-SAT Problem*. SAT 2006. pp. 252-265

² Joao Marques-Silva, Jordi Planes. *On Using Unsatisfiability for Solving Maximum Satisfiability*. CoRR abs/0712.1097. 2007

³ Joao Marques-Silva, Vasco M. Manquinho. *Towards More Effective Unsatisfiability-Based Maximum Satisfiability Algorithms*. SAT 2008. pp. 225-230

⁴ Carlos Ansótegui, Maria Luisa Bonet, Jordi Levy. *Solving (Weighted) Partial MaxSAT through Satisfiability Testing*. SAT 2009. pp. 427-440

⁵ Vasco M. Manquinho, Joao Marques Silva, Jordi Planes. *Algorithms for Weighted Boolean Optimization*. SAT 2009. pp. 495-508

(continued from previous page)

```
>>> 2
>>> print(fm.model)
[-1, -2, 3]
```

Module details

class `examples.fm.FM(formula, enc=0, solver='m22', verbose=1)`

A non-incremental implementation of the FM (Fu&Malik, or WMSU1) algorithm. The algorithm (see details in [Page 44, 5](#)) is *core-guided*, i.e. it solves maximum satisfiability with a series of unsatisfiability oracle calls, each producing an unsatisfiable core. The clauses involved in an unsatisfiable core are *relaxed* and a new `AtMost1` constraint on the corresponding *relaxation variables* is added to the formula. The process gets a bit more sophisticated in the case of weighted formulas because of the *clause weight splitting* technique.

The constructor of `FM` objects receives a target *WCNF* MaxSAT formula, an identifier of the cardinality encoding to use, a SAT solver name, and a verbosity level. Note that the algorithm uses the *pairwise* (see [card. EncType](#)) cardinality encoding by default, while the default SAT solver is MiniSat22 (referred to as 'm22', see [SolverNames](#) for details). The default verbosity level is 1.

Parameters

- **formula** (*WCNF*) – input MaxSAT formula
- **enc** (*int*) – cardinality encoding to use
- **solver** (*str*) – name of SAT solver
- **verbose** (*int*) – verbosity level

`_compute()`

This method implements WMSU1 algorithm. The method is essentially a loop, which at each iteration calls the SAT oracle to decide whether the working formula is satisfiable. If it is, the method derives a model (stored in variable `self.model`) and returns. Otherwise, a new unsatisfiable core of the formula is extracted and processed (see [treat_core\(\)](#)), and the algorithm proceeds.

`compute()`

Compute a MaxSAT solution. First, the method checks whether or not the set of hard clauses is satisfiable. If not, the method returns `False`. Otherwise, add soft clauses to the oracle and call the MaxSAT algorithm (see [_compute\(\)](#)).

Note that the soft clauses are added to the oracles after being augmented with additional *selector* literals. The selectors literals are then used as *assumptions* when calling the SAT oracle and are needed for extracting unsatisfiable cores.

`delete()`

Explicit destructor of the internal SAT oracle.

`init(with_soft=True)`

The method for the SAT oracle initialization. Since the oracle is used non-incrementally, it is reinitialized at every iteration of the MaxSAT algorithm (see [reinit\(\)](#)). An input parameter `with_soft` (`False` by default) regulates whether or not the formula's soft clauses are copied to the oracle.

Parameters `with_soft` (*bool*) – copy formula's soft clauses to the oracle or not

`oracle_time()`

Method for calculating and reporting the total SAT solving time.

reinit()

This method calls `delete()` and `init()` to reinitialize the internal SAT oracle. This is done at every iteration of the MaxSAT algorithm.

relax_core()

Relax and bound the core.

After unsatisfiable core splitting, this method is called. If the core contains only one clause, i.e. this clause cannot be satisfied together with the hard clauses of the formula, the formula gets augmented with the negation of the clause (see `remove_unit_core()`).

Otherwise (if the core contains more than one clause), every clause c of the core is *relaxed*. This means a new *relaxation literal* is added to the clause, i.e. $c \leftarrow c \vee r$, where r is a fresh (unused) relaxation variable. After the clauses get relaxed, a new cardinality encoding is added to the formula enforcing the sum of the new relaxation variables to be not greater than 1, $\sum_{c \in \phi} r \leq 1$, where ϕ denotes the unsatisfiable core.

remove_unit_core()

If an unsatisfiable core contains only one clause c , this method is invoked to add a bunch of new unit size hard clauses. As a result, the SAT oracle gets unit clauses $(\neg l)$ for all literals l in clause c .

split_core(minw)

Split clauses in the core whenever necessary.

Given a list of soft clauses in an unsatisfiable core, the method is used for splitting clauses whose weights are greater than the minimum weight of the core, i.e. the `minw` value computed in `treat_core()`. Each clause $(c \vee \neg s, w)$, s.t. $w > \text{minw}$ and s is its selector literal, is split into clauses (1) clause $(c \vee \neg s, \text{minw})$ and (2) a residual clause $(c \vee \neg s', w - \text{minw})$. Note that the residual clause has a fresh selector literal s' different from s .

Parameters `minw` (*int*) – minimum weight of the core

treat_core()

Now that the previous SAT call returned UNSAT, a new unsatisfiable core should be extracted and relaxed. Core extraction is done through a call to the `pysat.solvers.Solver.get_core()` method, which returns a subset of the selector literals deemed responsible for unsatisfiability.

After the core is extracted, its *minimum weight* `minw` is computed, i.e. it is the minimum weight among the weights of all soft clauses involved in the core (see [Page 44, 5](#)). Note that the cost of the MaxSAT solution is incremented by `minw`.

Clauses that have weight larger than `minw` are split (see `split_core()`). Afterwards, all clauses of the unsatisfiable core are relaxed (see `relax_core()`).

1.2.2 Hard formula generator (`pysat.examples.genhard`)

List of classes

<code>CB</code>	Mutilated chessboard principle (CB).
<code>GT</code>	Generator of ordering (or <i>greater than</i> , GT) principle formulas.
<code>PAR</code>	Generator of the parity principle (PAR) formulas.
<code>PHP</code>	Generator of k pigeonhole principle (k -PHP) formulas.

Module description

This module is designed to provide a few examples illustrating how PySAT can be used for encoding practical problems into CNF formulas. These include combinatorial principles that are widely studied from the propositional proof complexity perspective. Namely, encodings for the following principles are implemented: *pigeonhole principle* (*PHP*)¹, *ordering (greater-than) principle* (*GT*)², *mutilated chessboard principle* (*CB*)³, and *parity principle* (*PAR*)⁴.

The module can be used as an executable (the list of available command-line options can be shown using `genhard.py -h`) in the following way

```
$ genhard.py -t php -n 3 -v
c PHP formula for 4 pigeons and 3 holes
c (pigeon, hole) pair: (1, 1); bool var: 1
c (pigeon, hole) pair: (1, 2); bool var: 2
c (pigeon, hole) pair: (1, 3); bool var: 3
c (pigeon, hole) pair: (2, 1); bool var: 4
c (pigeon, hole) pair: (2, 2); bool var: 5
c (pigeon, hole) pair: (2, 3); bool var: 6
c (pigeon, hole) pair: (3, 1); bool var: 7
c (pigeon, hole) pair: (3, 2); bool var: 8
c (pigeon, hole) pair: (3, 3); bool var: 9
c (pigeon, hole) pair: (4, 1); bool var: 10
c (pigeon, hole) pair: (4, 2); bool var: 11
c (pigeon, hole) pair: (4, 3); bool var: 12
p cnf 12 22
1 2 3 0
4 5 6 0
7 8 9 0
10 11 12 0
-1 -4 0
-1 -7 0
-1 -10 0
-4 -7 0
-4 -10 0
-7 -10 0
-2 -5 0
-2 -8 0
-2 -11 0
-5 -8 0
-5 -11 0
-8 -11 0
-3 -6 0
-3 -9 0
-3 -12 0
-6 -9 0
-6 -12 0
-9 -12 0
```

Alternatively, each of the considered problem encoders can be accessed with the use of the standard `import` interface of Python, e.g.

¹ Stephen A. Cook, Robert A. Reckhow. *The Relative Efficiency of Propositional Proof Systems*. J. Symb. Log. 44(1). 1979. pp. 36-50

² Balakrishnan Krishnamurthy. *Short Proofs for Tricky Formulas*. Acta Informatica 22(3). 1985. pp. 253-275

³ Michael Alekhnovich. *Mutilated Chessboard Problem Is Exponentially Hard For Resolution*. Theor. Comput. Sci. 310(1-3). 2004. pp. 513-525

⁴ Miklós Ajtai. *Parity And The Pigeonhole Principle*. Feasible Mathematics. 1990. pp. 1-24

```
>>> from pysat.examples.genhard import PHP
>>>
>>> cnf = PHP(3)
>>> print(cnf.nv, len(cnf.clauses))
12 22
```

Given this example, observe that classes `PHP`, `GT`, `CB`, and `PAR` inherit from class `pysat.formula.CNF` and, thus, their corresponding clauses can be accessed through variable `.clauses`.

Module details

class `examples.genhard.CB`(*size*, *exhaustive=False*, *topv=0*, *verb=False*)

Mutilated chessboard principle (CB). Given an integer n , the principle states that it is impossible to cover a chessboard of size $2n \cdot 2n$ by domino tiles if two diagonally opposite corners of the chessboard are removed.

Note that the chessboard has $4n^2 - 2$ cells. Introduce a Boolean variable x_{ij} for $i, j \in [4n^2 - 2]$ s.t. cells i and j are adjacent (no variables are introduced for pairs of non-adjacent cells). CB formulas comprise clauses (1) $(\neg x_{ji} \vee \neg x_{ki})$ for every $i, j \neq k$ meaning that no more than one adjacent cell can be paired with the current one; and (2) $(\bigvee_{j \in \text{Adj}(i)} x_{ij}) \forall i$ enforcing that every cell i should be paired with at least one adjacent cell.

Clearly, since the two diagonal corners are removed, the formula is unsatisfiable. Also note the following. Assuming that the number of black cells is larger than the number of the white ones, CB formulas are unsatisfiable even if only a half of the formula is present, e.g. when `AtMost1` constraints are formulated only for the white cells while the `AtLeast1` constraints are formulated only for the black cells. Depending on the value of parameter `exhaustive` the encoder applies the *complete* or *partial* formulation of the problem.

Mutilated chessboard principle is known to be hard for resolution⁵.

Parameters

- **size** (*int*) – problem size (n)
- **exhaustive** (*bool*) – encode the problem exhaustively
- **topv** (*int*) – current top variable identifier
- **verb** (*bool*) – defines whether or not the encoder is verbose

Returns object of class `pysat.formula.CNF`.

class `examples.genhard.GT`(*size*, *topv=0*, *verb=False*)

Generator of ordering (or *greater than*, GT) principle formulas. Given an integer parameter n , the principle states that any partial order on the set $\{1, 2, \dots, n\}$ must have a maximal element.

Assume variable x_{ij} , for $i, j \in [n], i \neq j$, denotes the fact that $i \succ j$. Clauses $(\neg x_{ij} \vee \neg x_{ji})$ and $(\neg x_{ij} \vee \neg x_{jk} \vee x_{ik})$ ensure that the relation \succ is anti-symmetric and transitive. As a result, \succ is a partial order on $[n]$. The additional requirement that each element i has a successor in $[n] \setminus \{i\}$ represented a clause $(\bigvee_{j \neq i} x_{ji})$ makes the formula unsatisfiable.

GT formulas were originally conjectured⁷ to be hard for resolution. However,⁵ proved the existence of a polynomial size resolution refutation for GT formulas.

Parameters

- **size** (*int*) – number of elements (n)
- **topv** (*int*) – current top variable identifier
- **verb** (*bool*) – defines whether or not the encoder is verbose

⁵ Gunnar Stålmarck. *Short Resolution Proofs for a Sequence of Tricky Formulas*. Acta Informatica. 33(3). 1996. pp. 277-280

Returns object of class `pysat.formula.CNF`.

class `examples.genhard.PAR(size, topv=0, verb=False)`

Generator of the parity principle (PAR) formulas. Given an integer parameter n , the principle states that no graph on $2n + 1$ nodes consists of a complete perfect matching.

The encoding of the parity principle uses $\binom{2n+1}{2}$ variables $x_{ij}, i \neq j$. If variable x_{ij} is *true*, then there is an edge between nodes i and j . The formula consists of the following clauses: $(\bigvee_{j \neq i} x_{ij})$ for every $i \in [2n + 1]$, and $(\neg x_{ij} \vee \neg x_{kj})$ for all distinct $i, j, k \in [2n + 1]$.

The parity principle is known to be hard for resolution^{Page 47, 4}.

Parameters

- **size** (*int*) – problem size (n)
- **topv** (*int*) – current top variable identifier
- **verb** (*bool*) – defines whether or not the encoder is verbose

Returns object of class `pysat.formula.CNF`.

class `examples.genhard.PHP(nof_holes, kval=1, topv=0, verb=False)`

Generator of k pigeonhole principle (k -PHP) formulas. Given integer parameters m and k , the k pigeonhole principle states that if $k \cdot m + 1$ pigeons are distributed by m holes, then at least one hole contains more than k pigeons.

Note that if k is 1, the principle degenerates to the formulation of the original pigeonhole principle stating that $m + 1$ pigeons cannot be distributed by m holes.

Assume that a Boolean variable x_{ij} encodes that pigeon i resides in hole j . Then a PHP formula can be seen as a conjunction: $\bigwedge_{i=1}^{k \cdot m + 1} \text{AtLeast1}(x_{i1}, \dots, x_{im}) \wedge \bigwedge_{j=1}^m \text{AtMost}k(x_{1j}, \dots, x_{k \cdot m + 1, j})$. Here each **AtLeast1** constraint forces every pigeon to be placed into at least one hole while each **AtMost** k constraint allows the corresponding hole to have at most k pigeons. The overall PHP formulas are unsatisfiable.

PHP formulas are well-known⁶ to be hard for resolution.

Parameters

- **nof_holes** (*int*) – number of holes (n)
- **kval** (*int*) – multiplier k
- **topv** (*int*) – current top variable identifier
- **verb** (*bool*) – defines whether or not the encoder is verbose

Returns object of class `pysat.formula.CNF`.

1.2.3 Minimum/minimal hitting set solver (`pysat.examples.hitman`)

List of classes

Hitman

A cardinality-/subset-minimal hitting set enumerator.

⁶ Armin Haken. *The Intractability of Resolution*. Theor. Comput. Sci. 39. 1985. pp. 297-308

Module description

A SAT-based implementation of an implicit minimal hitting set¹ enumerator. The implementation is capable of computing/enumerating cardinality- and subset-minimal hitting sets of a given set of sets. Cardinality-minimal hitting set enumeration can be seen as ordered (sorted by size) subset-minimal hitting enumeration.

The minimal hitting set problem is trivially formulated as a MaxSAT formula in WCNF, as follows. Assume $E = \{e_1, \dots, e_n\}$ to be a universe of elements. Also assume there are k sets to hit: $s_i = \{e_{i,1}, \dots, e_{i,j_i}\}$ s.t. $e_{i,l} \in E$. Every set $s_i = \{e_{i,1}, \dots, e_{i,j_i}\}$ is translated into a hard clause $(e_{i,1} \vee \dots \vee e_{i,j_i})$. This results in the set of hard clauses having size k . The set of soft clauses comprises unit clauses of the form $(\neg e_j)$ s.t. $e_j \in E$, each having weight 1.

Taking into account this problem formulation as MaxSAT, ordered hitting enumeration is done with the use of the state-of-the-art MaxSAT solver called *RC2*^{2,3,4} while unordered hitting set enumeration is done through the *minimal correction subset* (MCS) enumeration, e.g. using the *LBX*⁵ or *MCS1s*-like⁶ MCS enumerators.

Hitman supports hitting set enumeration in the *implicit* manner, i.e. when sets to hit can be added on the fly as well as hitting sets can be blocked on demand.

An example usage of *Hitman* through the Python import interface is shown below. Here we target unordered subset-minimal hitting set enumeration.

```
>>> from pysat.examples.hitman import Hitman
>>>
>>> h = Hitman(solver='m22', htype='lbx')
>>> # adding sets to hit
>>> h.hit([1, 2, 3])
>>> h.hit([1, 4])
>>> h.hit([5, 6, 7])
>>>
>>> h.get()
[1, 5]
>>>
>>> h.block([1, 5])
>>>
>>> h.get()
[2, 4, 5]
>>>
>>> h.delete()
```

Enumerating cardinality-minimal hitting sets can be done as follows:

```
>>> from pysat.examples.hitman import Hitman
>>>
>>> sets = [[1, 2, 3], [1, 4], [5, 6, 7]]
>>> with Hitman(bootstrap_with=sets, htype='sorted') as hitman:
...     for hs in hitman.enumerate():
...         print(hs)
...
```

(continues on next page)

¹ Erick Moreno-Centeno, Richard M. Karp. *The Implicit Hitting Set Approach to Solve Combinatorial Optimization Problems with an Application to Multigenome Alignment*. Operations Research 61(2). 2013. pp. 453-468

² António Morgado, Carmine Dodaro, Joao Marques-Silva. *Core-Guided MaxSAT with Soft Cardinality Constraints*. CP 2014. pp. 564-573

³ António Morgado, Alexey Ignatiev, Joao Marques-Silva. *MSCG: Robust Core-Guided MaxSAT Solving*. JSAT 9. 2014. pp. 129-134

⁴ Alexey Ignatiev, António Morgado, Joao Marques-Silva. *RC2: a Python-based MaxSAT Solver*. MaxSAT Evaluation 2018. p. 22

⁵ Carlos Mencía, Alessandro Previti, Joao Marques-Silva. *Literal-Based MCS Extraction*. IJCAI. 2015. pp. 1973-1979

⁶ Joao Marques-Silva, Federico Heras, Mikolás Janota, Alessandro Previti, Anton Belov. *On Computing Minimal Correction Subsets*. IJCAI. 2013. pp. 615-622

(continued from previous page)

```
[1, 5]
[1, 6]
[1, 7]
[3, 4, 7]
[2, 4, 7]
[3, 4, 6]
[3, 4, 5]
[2, 4, 6]
[2, 4, 5]
```

Finally, implicit hitting set enumeration can be used in practical problem solving. As an example, let us show the basic flow of a MaxHS-like⁷ algorithm for MaxSAT:

```
>>> from pysat.examples.hitman import Hitman
>>> from pysat.solvers import Solver
>>>
>>> hitman = Hitman(htype='sorted')
>>> oracle = Solver()
>>>
>>> # here we assume that the SAT oracle
>>> # is initialized with a MaxSAT formula,
>>> # whose soft clauses are extended with
>>> # selector literals stored in "sels"
>>> while True:
...     hs = hitman.get() # hitting the set of unsatisfiable cores
...     ts = set(sels).difference(set(hs)) # soft clauses to try
...
...     if oracle.solve(assumptions=ts):
...         print('s OPTIMUM FOUND')
...         print('o', len(hs))
...         break
...     else:
...         core = oracle.get_core()
...         hitman.hit(core)
```

Module details

class `examples.hitman.Atom(obj, sign=True)`

Atoms are elementary (signed) objects necessary when dealing with hitting sets subject to hard constraints.

class `examples.hitman.Hitman(bootstrap_with=[], weights=None, subject_to=[], solver='g3', htype='sorted', mxs_adapt=False, mxs_exhaust=False, mxs_minz=False, mxs_trim=0, mcs_usecld=False)`

A cardinality-/subset-minimal hitting set enumerator. The enumerator can be set up to use either a MaxSAT solver [RC2](#) or an MCS enumerator (either [LBX](#) or [MCS1s](#)). In the former case, the hitting sets enumerated are ordered by size (smallest size hitting sets are computed first), i.e. *sorted*. In the latter case, subset-minimal hitting are enumerated in an arbitrary order, i.e. *unsorted*.

This is handled with the use of parameter `htype`, which is set to be `'sorted'` by default. The MaxSAT-based enumerator can be chosen by setting `htype` to one of the following values: `'maxsat'`, `'mxs'`, or `'rc2'`.

⁷ Jessica Davies, Fahiem Bacchus. *Solving MAXSAT by Solving a Sequence of Simpler SAT Instances*. CP 2011. pp. 225-239

Alternatively, by setting it to 'mcs' or 'lbx', a user can enforce using the [LBX](#) MCS enumerator. If `htype` is set to 'mcs1s', the [MCS1s](#) enumerator is used.

In either case, an underlying problem solver can use a SAT oracle specified as an input parameter `solver`. The default SAT solver is Glucose3 (specified as `g3`, see [SolverNames](#) for details).

Objects of class [Hitman](#) can be bootstrapped with an iterable of iterables, e.g. a list of lists. This is handled using the `bootstrap_with` parameter. Each set to hit can comprise elements of any type, e.g. integers, strings or objects of any Python class, as well as their combinations. The bootstrapping phase is done in `init()`.

Another optional parameter `subject_to` can be used to specify arbitrary hard constraints that must be respected when computing hitting sets of the given sets. Note that `subject_to` should be an iterable containing pure clauses and/or native `AtMostK` constraints. Note that native cardinality constraints supported only by MiniCard-like solvers. Finally, note that these hard constraints must be defined over the set of signed atomic objects, i.e. instances of class [Atom](#).

A few other optional parameters include the possible options for RC2 as well as for LBX- and MCS1s-like MCS enumerators that control the behaviour of the underlying solvers.

Parameters

- **`bootstrap_with`** (*iterable(iterable(obj))*) – input set of sets to hit
- **`weights`** (*dict(obj)*) – a mapping from objects to their weights (if weighted)
- **`subject_to`** (*iterable(iterable([Atom](#)))*) – hard constraints (either clauses or native `AtMostK` constraints)
- **`solver`** (*str*) – name of SAT solver
- **`htype`** (*str*) – enumerator type
- **`mxs_adapt`** (*bool*) – detect and process `AtMost1` constraints in RC2
- **`mxs_exhaust`** (*bool*) – apply unsatisfiable core exhaustion in RC2
- **`mxs_minz`** (*bool*) – apply heuristic core minimization in RC2
- **`mxs_trim`** (*int*) – trim unsatisfiable cores at most this number of times
- **`mcs_useclid`** (*bool*) – use clause-D heuristic in the MCS enumerator

`add_hard`(*clause, weights=None*)

Add a hard constraint, which can be either a pure clause or an `AtMostK` constraint.

Note that an optional parameter that can be passed to this method is `weights`, which contains a mapping the objects under question into weights. Also note that the weight of an object must not change from one call of `hit()` to another.

Parameters

- **`clause`** (*iterable(obj)*) – hard constraint (either a clause or a native `AtMostK` constraint)
- **`weights`** (*dict(obj)*) – a mapping from objects to weights

`block`(*to_block, weights=None*)

The method serves for imposing a constraint forbidding the hitting set solver to compute a given hitting set. Each set to block is encoded as a hard clause in the MaxSAT problem formulation, which is then added to the underlying oracle.

Note that an optional parameter that can be passed to this method is `weights`, which contains a mapping the objects under question into weights. Also note that the weight of an object must not change from one call of `hit()` to another.

Parameters

- **to_block** (*iterable(obj)*) – a set to block
- **weights** (*dict(obj)*) – a mapping from objects to weights

delete()

Explicit destructor of the internal hitting set oracle.

enumerate()

The method can be used as a simple iterator computing and blocking the hitting sets on the fly. It essentially calls *get()* followed by *block()*. Each hitting set is reported as a list of objects in the original problem domain, i.e. it is mapped back from the solutions over Boolean variables computed by the underlying oracle.

Return type list(obj)

get()

This method computes and returns a hitting set. The hitting set is obtained using the underlying oracle operating the MaxSAT problem formulation. The computed solution is mapped back to objects of the problem domain.

Return type list(obj)

hit(to_hit, weights=None)

This method adds a new set to hit to the hitting set solver. This is done by translating the input iterable of objects into a list of Boolean variables in the MaxSAT problem formulation.

Note that an optional parameter that can be passed to this method is **weights**, which contains a mapping the objects under question into weights. Also note that the weight of an object must not change from one call of *hit()* to another.

Parameters

- **to_hit** (*iterable(obj)*) – a new set to hit
- **weights** (*dict(obj)*) – a mapping from objects to weights

init(bootstrap_with, weights=None, subject_to=[])

This method serves for initializing the hitting set solver with a given list of sets to hit. Concretely, the hitting set problem is encoded into partial MaxSAT as outlined above, which is then fed either to a MaxSAT solver or an MCS enumerator.

An additional optional parameter is **weights**, which can be used to specify non-unit weights for the target objects in the sets to hit. This only works if 'sorted' enumeration of hitting sets is applied.

Another optional parameter is available, namely, **subject_to**. It can be used to specify arbitrary hard constraints that must be respected when computing hitting sets of the given sets. Note that **subject_to** should be an iterable containing pure clauses and/or native AtMostK constraints. Finally, note that these hard constraints must be defined over the set of signed atomic objects, i.e. instances of class *Atom*.

Parameters

- **bootstrap_with** (*iterable(iterable(obj))*) – input set of sets to hit
- **weights** (*dict(obj)*) – weights of the objects in case the problem is weighted
- **subject_to** (*iterable(iterable(Atom))*) – hard constraints (either clauses or native AtMostK constraints)

oracle_time()

Report the total SAT solving time.

1.2.4 LBX-like MCS enumerator (`pysat.examples.lbx`)

List of classes

<i>LBX</i>	LBX-like algorithm for computing MCSes.
------------	---

Module description

This module implements a prototype of the LBX algorithm for the computation of a *minimal correction subset* (MCS) and/or MCS enumeration. The LBX abbreviation stands for *literal-based MCS extraction* algorithm, which was proposed in¹. Note that this prototype does not follow the original low-level implementation of the corresponding MCS extractor available [online](#) (compared to our prototype, the low-level implementation has a number of additional heuristics used). However, it implements the LBX algorithm for partial MaxSAT formulas, as described in¹.

The implementation can be used as an executable (the list of available command-line options can be shown using `lbx.py -h`) in the following way:

```
$ xzcat formula.wcnf.xz
p wcnf 3 6 4
1 1 0
1 2 0
1 3 0
4 -1 -2 0
4 -1 -3 0
4 -2 -3 0

$ lbx.py -d -e all -s glucose3 -vv formula.wcnf.xz
c MCS: 1 3 0
c cost: 2
c MCS: 2 3 0
c cost: 2
c MCS: 1 2 0
c cost: 2
c oracle time: 0.0002
```

Alternatively, the algorithm can be accessed and invoked through the standard `import` interface of Python, e.g.

```
>>> from pysat.examples.lbx import LBX
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF(from_file='formula.wcnf.xz')
>>>
>>> lbx = LBX(wcnf, use_cld=True, solver_name='g3')
>>> for mcs in lbx.enumerate():
...     lbx.block(mcs)
...     print(mcs)
[1, 3]
[2, 3]
[1, 2]
```

¹ Carlos Mencia, Alessandro Previti, Joao Marques-Silva. *Literal-Based MCS Extraction*. IJCAI 2015. pp. 1973-1979

Module details

class `examples.lbx.LBX`(*formula*, *use_cld=False*, *solver_name='m22'*, *use_timer=False*)

LBX-like algorithm for computing MCSes. Given an unsatisfiable partial CNF formula, i.e. formula in the [WCNF](#) format, this class can be used to compute a given number of MCSes of the formula. The implementation follows the LBX algorithm description in [Page 54, 1](#). It can use any SAT solver available in PySAT. Additionally, the “clause *D*” heuristic can be used when enumerating MCSes.

The default SAT solver to use is `m22` (see [SolverNames](#)). The “clause *D*” heuristic is disabled by default, i.e. `use_cld` is set to `False`. Internal SAT solver’s timer is also disabled by default, i.e. `use_timer` is `False`.

Parameters

- **formula** ([WCNF](#)) – unsatisfiable partial CNF formula
- **use_cld** (*bool*) – whether or not to use “clause *D*”
- **solver_name** (*str*) – SAT oracle name
- **use_timer** (*bool*) – whether or not to use SAT solver’s timer

`_compute()`

The main method of the class, which computes an MCS given its over-approximation. The over-approximation is defined by a model for the hard part of the formula obtained in [compute\(\)](#).

The method is essentially a simple loop going over all literals unsatisfied by the previous model, i.e. the literals of `self.setd` and checking which literals can be satisfied. This process can be seen a refinement of the over-approximation of the MCS. The algorithm follows the pseudo-code of the LBX algorithm presented in [Page 54, 1](#).

Additionally, if [LBX](#) was constructed with the requirement to make “clause *D*” calls, the method calls [do_cld_check\(\)](#) at every iteration of the loop using the literals of `self.setd` not yet checked, as the contents of “clause *D*”.

`_filter_satisfied(update_setd=False)`

This method extracts a model provided by the previous call to a SAT oracle and iterates over all soft clauses checking if each of is satisfied by the model. Satisfied clauses are marked accordingly while the literals of the unsatisfied clauses are kept in a list called `setd`, which is then used to refine the correction set (see [_compute\(\)](#), and [do_cld_check\(\)](#)).

Optional Boolean parameter `update_setd` enforces the method to update variable `self.setd`. If this parameter is set to `False`, the method only updates the list of satisfied clauses, which is an under-approximation of a *maximal satisfiable subset* (MSS).

Parameters `update_setd` (*bool*) – whether or not to update `setd`

`_map_extlit(l)`

Map an external variable to an internal one if necessary.

This method is used when new clauses are added to the formula incrementally, which may result in introducing new variables clashing with the previously used *clause selectors*. The method makes sure no clash occurs, i.e. it maps the original variables used in the new problem clauses to the newly introduced auxiliary variables (see [add_clause\(\)](#)).

Given an integer literal, a fresh literal is returned. The returned integer has the same sign as the input literal.

Parameters `l` (*int*) – literal to map

Return type `int`

_satisfied(*cl, model*)

Given a clause (as an iterable of integers) and an assignment (as a list of integers), this method checks whether or not the assignment satisfies the clause. This is done by a simple clause traversal. The method is invoked from `_filter_satisfied()`.

Parameters

- **cl** (*iterable(int)*) – a clause to check
- **model** (*list(int)*) – an assignment

Return type bool**add_clause**(*clause, soft=False*)

The method for adding a new hard or soft clause to the problem formula. Although the input formula is to be specified as an argument of the constructor of `LBX`, adding clauses may be helpful when *enumerating* MCSes of the formula. This way, the clauses are added incrementally, i.e. *on the fly*.

The clause to add can be any iterable over integer literals. The additional Boolean parameter `soft` can be set to `True` meaning the the clause being added is soft (note that parameter `soft` is set to `False` by default).

Also note that besides pure clauses, the method can also expect native cardinality constraints represented as a pair (`lits`, `bound`). Only hard cardinality constraints can be added.

Parameters

- **clause** (*iterable(int)*) – a clause to add
- **soft** (*bool*) – whether or not the clause is soft

block(*mcs*)

Block a (previously computed) MCS. The MCS should be given as an iterable of integers. Note that this method is not automatically invoked from `enumerate()` because a user may want to block some of the MCSes conditionally depending on the needs. For example, one may want to compute disjoint MCSes only in which case this standard blocking is not appropriate.

Parameters **mcs** (*iterable(int)*) – an MCS to block

compute(*enable=[]*)

Compute and return one solution. This method checks whether the hard part of the formula is satisfiable, i.e. an MCS can be extracted. If the formula is satisfiable, the model computed by the SAT call is used as an *over-approximation* of the MCS in the method `_compute()` invoked here, which implements the LBX algorithm.

An MCS is reported as a list of integers, each representing a soft clause index (the smallest index is 1).

An optional input parameter is `enable`, which represents a sequence (normally a list) of soft clause indices that a user would prefer to enable/satisfy. Note that this may result in an unsatisfiable oracle call, in which case `None` will be reported as solution. Also, the smallest clause index is assumed to be 1.

Parameters **enable** (*iterable(int)*) – a sequence of clause ids to enable

Return type list(int)

delete()

Explicit destructor of the internal SAT oracle.

do_cld_check(*cld*)

Do the “clause *D*” check. This method receives a list of literals, which serves a “clause *D*”², and checks whether the formula conjoined with *D* is satisfiable.

² Joao Marques-Silva, Federico Heras, Mikolas Janota, Alessandro Previti, Anton Belov. *On Computing Minimal Correction Subsets*. IJCAI 2013. pp. 615-622

If clause D cannot be satisfied together with the formula, then negations of all of its literals are backbones of the formula and the LBX algorithm can stop. Otherwise, the literals satisfied by the new model refine the MCS further.

Every time the method is called, a new fresh selector variable s is introduced, which augments the current clause D . The SAT oracle then checks if clause $(D \vee \neg s)$ can be satisfied together with the internal formula. The D clause is then disabled by adding a hard clause $(\neg s)$.

Parameters `cld` (`list(int)`) – clause D to check

enumerate()

This method iterates through MCSes enumerating them until the formula has no more MCSes. The method iteratively invokes `compute()`. Note that the method does not block the MCSes computed - this should be explicitly done by a user.

oracle_time()

Report the total SAT solving time.

1.2.5 LSU algorithm for MaxSAT (`pysat.examples.lsu`)

List of classes

<code>LSU</code>	Linear SAT-UNSAT algorithm for MaxSAT ¹ .
<code>LSUPlus</code>	LSU-like algorithm extended for <code>WCNFPlus</code> formulas (using Minicard).

Module description

The module implements a prototype of the known *LSU/LSUS*, e.g. *linear (search) SAT-UNSAT*, algorithm for MaxSAT, e.g. see¹. The implementation is improved by the use of the *iterative totalizer encoding*². The encoding is used in an incremental fashion, i.e. it is created once and reused as many times as the number of iterations the algorithm makes.

The implementation can be used as an executable (the list of available command-line options can be shown using `lsu.py -h`) in the following way:

```
$ xzcat formula.wcnf.xz
p wcnf 3 6 4
1 1 0
1 2 0
1 3 0
4 -1 -2 0
4 -1 -3 0
4 -2 -3 0

$ lsu.py -s glucose3 -m -v formula.wcnf.xz
c formula: 3 vars, 3 hard, 3 soft
o 2
s OPTIMUM FOUND
v -1 -2 3 0
c oracle time: 0.0000
```

¹ António Morgado, Federico Heras, Mark H. Liffiton, Jordi Planes, Joao Marques-Silva. *Iterative and core-guided MaxSAT solving: A survey and assessment*. Constraints 18(4). 2013. pp. 478-534

² Ruben Martins, Saurabh Joshi, Vasco M. Manquinho, Inês Lynce. *Incremental Cardinality Constraints for MaxSAT*. CP 2014. pp. 531-548

Alternatively, the algorithm can be accessed and invoked through the standard `import` interface of Python, e.g.

```
>>> from pysat.examples.lsu import LSU
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF(from_file='formula.wcnf.xz')
>>>
>>> lsu = LSU(wcnf, verbose=0)
>>> lsu.solve() # set of hard clauses should be satisfiable
True
>>> print(lsu.cost) # cost of MaxSAT solution should be 2
>>> 2
>>> print(lsu.model)
[-1, -2, 3]
```

Module details

class `examples.lsu.LSU`(*formula*, *solver*='g4', *expect_interrupt*=False, *verbose*=0)

Linear SAT-UNSAT algorithm for MaxSAT[?]. The algorithm can be seen as a series of satisfiability oracle calls refining an upper bound on the MaxSAT cost, followed by one unsatisfiability call, which stops the algorithm. The implementation encodes the sum of all selector literals using the *iterative totalizer encoding*^{Page 57, 2}. At every iteration, the upper bound on the cost is reduced and enforced by adding the corresponding unit size clause to the working formula. No clauses are removed during the execution of the algorithm. As a result, the SAT oracle is used incrementally.

Warning: At this point, *LSU* supports only **unweighted** problems.

The constructor receives an input *WCNF* formula, a name of the SAT solver to use (see *SolverNames* for details), and an integer verbosity level.

Parameters

- **formula** (*WCNF*) – input MaxSAT formula
- **solver** (*str*) – name of SAT solver
- **expect_interrupt** (*bool*) – whether or not an *interrupt()* call is expected
- **verbose** (*int*) – verbosity level

`_assert_lt(cost)`

The method enforces an upper bound on the cost of the MaxSAT solution. This is done by encoding the sum of all soft clause selectors with the use the iterative totalizer encoding, i.e. *ITotalizer*. Note that the sum is created once, at the beginning. Each of the following calls to this method only enforces the upper bound on the created sum by adding the corresponding unit size clause. Each such clause is added on the fly with no restart of the underlying SAT oracle.

Parameters *cost* (*int*) – the cost of the next MaxSAT solution is enforced to be *lower* than this current cost

`_get_model_cost(formula, model)`

Given a WCNF formula and a model, the method computes the MaxSAT cost of the model, i.e. the sum of weights of soft clauses that are unsatisfied by the model.

Parameters

- **formula** (*WCNF*) – an input MaxSAT formula
- **model** (*list(int)*) – a satisfying assignment

Return type int

_init(*formula*)

SAT oracle initialization. The method creates a new SAT oracle and feeds it with the formula's hard clauses. Afterwards, all soft clauses of the formula are augmented with selector literals and also added to the solver. The list of all introduced selectors is stored in variable `self.sels`.

Parameters **formula** (*WCNF*) – input MaxSAT formula

clear_interrupt()

Clears an interruption.

delete()

Explicit destructor of the internal SAT oracle and the *ITotalizer* object.

found_optimum()

Checks if the optimum solution was found in a prior call to *solve*().

Return type bool

get_model()

This method returns a model obtained during a prior satisfiability oracle call made in *solve*().

Return type list(int)

interrupt()

Interrupt the current execution of LSU's *solve*() method. Can be used to enforce time limits using timer objects. The interrupt must be cleared before running the LSU algorithm again (see *clear_interrupt*()).

oracle_time()

Method for calculating and reporting the total SAT solving time.

solve()

Computes a solution to the MaxSAT problem. The method implements the LSU/LSUS algorithm, i.e. it represents a loop, each iteration of which calls a SAT oracle on the working MaxSAT formula and refines the upper bound on the MaxSAT cost until the formula becomes unsatisfiable.

Returns True if the hard part of the MaxSAT formula is satisfiable, i.e. if there is a MaxSAT solution, and False otherwise.

Return type bool

class `examples.lsu.LSUPlus`(*formula*, *solver*, *expect_interrupt=False*, *verbose=0*)

LSU-like algorithm extended for *WCNFPlus* formulas (using Minicard).

Parameters

- **formula** (*WCNFPlus*) – input MaxSAT formula in WCNF+ format
- **expect_interrupt** (*bool*) – whether or not an *interrupt*() call is expected
- **verbose** (*int*) – verbosity level

_assert_lt(*cost*)

Overrides *_assert_lt* of *LSU* in order to use Minicard's native support for cardinality constraints

Parameters **cost** (*int*) – the cost of the next MaxSAT solution is enforced to be *lower* than this current cost

1.2.6 CLD-like MCS enumerator (`pysat.examples.mcsls`)

List of classes

<code>MCSls</code>	Algorithm BLS for computing MCSes, augmented with "clause <i>D</i> " calls.
--------------------	---

Module description

This module implements a prototype of a BLS- and CLD-like algorithm for the computation of a *minimal correction subset* (MCS) and/or MCS enumeration. More concretely, the implementation follows the *basic linear search* (BLS) for MCS extraction augmented with *clause D* (CLD) oracle calls. As a result, the algorithm is not an implementation of the BLS or CLD algorithms as described in¹ but a mixture of both. Note that the corresponding original low-level implementations of both can be found [online](#).

The implementation can be used as an executable (the list of available command-line options can be shown using `mcsls.py -h`) in the following way:

```
$ xzcat formula.wcnf.xz
p wcnf 3 6 4
1 1 0
1 2 0
1 3 0
4 -1 -2 0
4 -1 -3 0
4 -2 -3 0

$ mcsls.py -d -e all -s glucose3 -vv formula.wcnf.xz
c MCS: 1 3 0
c cost: 2
c MCS: 2 3 0
c cost: 2
c MCS: 1 2 0
c cost: 2
c oracle time: 0.0002
```

Alternatively, the algorithm can be accessed and invoked through the standard `import` interface of Python, e.g.

```
>>> from pysat.examples.mcsls import MCSls
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF(from_file='formula.wcnf.xz')
>>>
>>> mcsls = MCSls(wcnf, use_cld=True, solver_name='g3')
>>> for mcs in mcsls.enumerate():
...     mcsls.block(mcs)
...     print(mcs)
[1, 3]
[2, 3]
[1, 2]
```

¹ Joao Marques-Silva, Federico Heras, Mikolas Janota, Alessandro Previti, Anton Belov. *On Computing Minimal Correction Subsets*. IJCAI 2013. pp. 615-622

Module details

class `examples.mcs1s.MCS1s`(*formula*, *use_cld=False*, *solver_name='m22'*, *use_timer=False*)

Algorithm BLS for computing MCSes, augmented with “clause *D*” calls. Given an unsatisfiable partial CNF formula, i.e. formula in the [WCNF](#) format, this class can be used to compute a given number of MCSes of the formula. The implementation follows the description of the basic linear search (BLS) algorithm description in [Page 60, 1](#). It can use any SAT solver available in PySAT. Additionally, the “clause *D*” heuristic can be used when enumerating MCSes.

The default SAT solver to use is `m22` (see [SolverNames](#)). The “clause *D*” heuristic is disabled by default, i.e. `use_cld` is set to `False`. Internal SAT solver’s timer is also disabled by default, i.e. `use_timer` is `False`.

Parameters

- **formula** ([WCNF](#)) – unsatisfiable partial CNF formula
- **use_cld** (*bool*) – whether or not to use “clause *D*”
- **solver_name** (*str*) – SAT oracle name
- **use_timer** (*bool*) – whether or not to use SAT solver’s timer

`_compute()`

The main method of the class, which computes an MCS given its over-approximation. The over-approximation is defined by a model for the hard part of the formula obtained in `_overapprox()` (the corresponding oracle is made in `compute()`).

The method is essentially a simple loop going over all literals unsatisfied by the previous model, i.e. the literals of `self.setd` and checking which literals can be satisfied. This process can be seen a refinement of the over-approximation of the MCS. The algorithm follows the pseudo-code of the BLS algorithm presented in [Page 60, 1](#).

Additionally, if `MCS1s` was constructed with the requirement to make “clause *D*” calls, the method calls `do_cld_check()` at every iteration of the loop using the literals of `self.setd` not yet checked, as the contents of “clause *D*”.

`_map_extlit(l)`

Map an external variable to an internal one if necessary.

This method is used when new clauses are added to the formula incrementally, which may result in introducing new variables clashing with the previously used *clause selectors*. The method makes sure no clash occurs, i.e. it maps the original variables used in the new problem clauses to the newly introduced auxiliary variables (see `add_clause()`).

Given an integer literal, a fresh literal is returned. The returned integer has the same sign as the input literal.

Parameters `l` (*int*) – literal to map

Return type `int`

`_overapprox()`

The method extracts a model corresponding to an over-approximation of an MCS, i.e. it is the model of the hard part of the formula (the corresponding oracle call is made in `compute()`).

Here, the set of selectors is divided into two parts: `self.ss_assumps`, which is an under-approximation of an MSS (maximal satisfiable subset) and `self.setd`, which is an over-approximation of the target MCS. Both will be further refined in `_compute()`.

`add_clause`(*clause*, *soft=False*)

The method for adding a new hard or soft clause to the problem formula. Although the input formula is to

be specified as an argument of the constructor of `MCSIs`, adding clauses may be helpful when *enumerating* MCSes of the formula. This way, the clauses are added incrementally, i.e. *on the fly*.

The clause to add can be any iterable over integer literals. The additional Boolean parameter `soft` can be set to `True` meaning the the clause being added is soft (note that parameter `soft` is set to `False` by default).

Also note that besides pure clauses, the method can also expect native cardinality constraints represented as a pair `(lits, bound)`. Only hard cardinality constraints can be added.

Parameters

- **clause** (*iterable(int)*) – a clause to add
- **soft** (*bool*) – whether or not the clause is soft

`block(mcs)`

Block a (previously computed) MCS. The MCS should be given as an iterable of integers. Note that this method is not automatically invoked from `enumerate()` because a user may want to block some of the MCSes conditionally depending on the needs. For example, one may want to compute disjoint MCSes only in which case this standard blocking is not appropriate.

Parameters `mcs` (*iterable(int)*) – an MCS to block

`compute(enable=[])`

Compute and return one solution. This method checks whether the hard part of the formula is satisfiable, i.e. an MCS can be extracted. If the formula is satisfiable, the model computed by the SAT call is used as an *over-approximation* of the MCS in the method `_compute()` invoked here, which implements the BLS

An MCS is reported as a list of integers, each representing a soft clause index (the smallest index is 1).

An optional input parameter is `enable`, which represents a sequence (normally a list) of soft clause indices that a user would prefer to enable/satisfy. Note that this may result in an unsatisfiable oracle call, in which case `None` will be reported as solution. Also, the smallest clause index is assumed to be 1.

Parameters `enable` (*iterable(int)*) – a sequence of clause ids to enable

Return type `list(int)`

`delete()`

Explicit destructor of the internal SAT oracle.

`do_cld_check(cld)`

Do the “clause *D*” check. This method receives a list of literals, which serves a “clause *D*”^{Page 60, 1}, and checks whether the formula conjoined with *D* is satisfiable.

If clause *D* cannot be satisfied together with the formula, then negations of all of its literals are backbones of the formula and the MCSIs algorithm can stop. Otherwise, the literals satisfied by the new model refine the MCS further.

Every time the method is called, a new fresh selector variable *s* is introduced, which augments the current clause *D*. The SAT oracle then checks if clause $(D \vee \neg s)$ can be satisfied together with the internal formula. The *D* clause is then disabled by adding a hard clause $(\neg s)$.

Parameters `cld` (*list(int)*) – clause *D* to check

`enumerate()`

This method iterates through MCSes enumerating them until the formula has no more MCSes. The method iteratively invokes `compute()`. Note that the method does not block the MCSes computed - this should be explicitly done by a user.

`oracle_time()`

Report the total SAT solving time.

1.2.7 An iterative model enumerator (`pysat.examples.models`)

List of classes

<code>enumerate_models</code>	Enumeration procedure.
-------------------------------	------------------------

Module description

The module implements a simple iterative enumeration of a given number of models of [CNF](#) or CNFPlus formula. In the latter case, only Minicard can be used as a SAT solver. The module aims at illustrating how one can work with model computation and enumeration.

The implementation facilitates the simplest use of a SAT oracle from the *command line*. If one deals with the enumeration task from a Python script, it is more convenient to exploit the internal model enumeration of the `pysat.solvers` module. Concretely, see `pysat.solvers.Solver.enum_models()`.

```
$ cat formula.cnf
p cnf 4 4
-1 2 0
-1 3 0
-2 4 0
3 -4 0

$ models.py -e all -s glucose3 formula.cnf
v -1 -2 +3 -4 0
v +1 +2 -3 +4 0
c nof models: 2
c accum time: 0.00s
c mean  time: 0.00s
```

Module details

`examples.models.enumerate_models(formula, to_enum, solver, warm=False)`

Enumeration procedure. It represents a loop iterating over satisfying assignment for a given formula until either all or a given number of them is enumerated.

Parameters

- **formula** ([CNFPlus](#)) – input WCNF formula
- **to_enum** (*int* or *'all'*) – number of models to compute
- **solver** (*str*) – name of SAT solver
- **warm** (*bool*) – warm start flag

1.2.8 A deletion-based MUS extractor (`pysat.examples.musx`)

List of classes

*MUSX*MUS eXtractor using the deletion-based algorithm.

Module description

This module implements a deletion-based algorithm¹ for extracting a *minimal unsatisfiable subset (MUS)* of a given (unsatisfiable) CNF formula. This simplistic implementation can deal with *plain* and *partial* CNF formulas, e.g. formulas in the DIMACS CNF and WCNF formats.

The following extraction procedure is implemented:

```
# oracle: SAT solver (initialized)
# assump: full set of assumptions

i = 0

while i < len(assump):
    to_test = assump[:i] + assump[(i + 1):]
    if oracle.solve(assumptions=to_test):
        i += 1
    else:
        assump = to_test

return assump
```

The implementation can be used as an executable (the list of available command-line options can be shown using `musx.py -h`) in the following way:

```
$ cat formula.wcnf
p wcnf 3 6 4
1 1 0
1 2 0
1 3 0
4 -1 -2 0
4 -1 -3 0
4 -2 -3 0

$ musx.py -s glucose3 -vv formula.wcnf
c MUS approx: 1 2 0
c testing clid: 0 -> sat (keeping 0)
c testing clid: 1 -> sat (keeping 1)
c nof soft: 3
c MUS size: 2
v 1 2 0
c oracle time: 0.0001
```

Alternatively, the algorithm can be accessed and invoked through the standard `import` interface of Python, e.g.

¹ Joao Marques-Silva. *Minimal Unsatisfiability: Models, Algorithms and Applications*. ISMVL 2010. pp. 9-14

```

>>> from pysat.examples.musx import MUSX
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF(from_file='formula.wcnf')
>>>
>>> musx = MUSX(wcnf, verbosity=0)
>>> musx.compute() # compute a minimally unsatisfiable set of clauses
[1, 2]

```

Note that the implementation is able to compute only one MUS (MUS enumeration is not supported).

Module details

class `examples.musx.MUSX(formula, solver='m22', verbosity=1)`

MUS eXtractor using the deletion-based algorithm. The algorithm is described in [Page 64, 1](#) (also see the module description above). Essentially, the algorithm can be seen as an iterative process, which tries to remove one soft clause at a time and check whether the remaining set of soft clauses is still unsatisfiable together with the hard clauses.

The constructor of `MUSX` objects receives a target `WCNF` formula, a SAT solver name, and a verbosity level. Note that the default SAT solver is MiniSat22 (referred to as 'm22', see [SolverNames](#) for details). The default verbosity level is 1.

Parameters

- **formula** (`WCNF`) – input WCNF formula
- **solver** (`str`) – name of SAT solver
- **verbosity** (`int`) – verbosity level

`_compute(approx)`

Deletion-based MUS extraction. Given an over-approximation of an MUS, i.e. an unsatisfiable core previously returned by a SAT oracle, the method represents a loop, which at each iteration removes a clause from the core and checks whether the remaining clauses of the approximation are unsatisfiable together with the hard clauses.

Soft clauses are (de)activated using the standard MiniSat-like assumptions interface². Each soft clause c is augmented with a selector literal s , e.g. $(c) \leftarrow (c \vee \neg s)$. As a result, clause c can be activated by assuming literal s . The over-approximation provided as an input is specified as a list of selector literals for clauses in the unsatisfiable core.

Parameters `approx` (`list(int)`) – an over-approximation of an MUS

Note that the method does not return. Instead, after its execution, the input over-approximation is refined and contains an MUS.

`compute()`

This is the main method of the `MUSX` class. It computes a set of soft clauses belonging to an MUS of the input formula. First, the method checks whether the formula is satisfiable. If it is, nothing else is done. Otherwise, an *unsatisfiable core* of the formula is extracted, which is later used as an over-approximation of an MUS refined in `_compute()`.

`delete()`

Explicit destructor of the internal SAT oracle.

² Niklas Eén, Niklas Sörensson. *Temporal induction by incremental SAT solving*. Electr. Notes Theor. Comput. Sci. 89(4). 2003. pp. 543-560

`oracle_time()`

Method for calculating and reporting the total SAT solving time.

1.2.9 OptUx optimal MUS enumerator (`pysat.examples.optux`)

List of classes

<code>OptUx</code>	A simple Python version of the implicit hitting set based optimal MUS extractor and enumerator.
--------------------	---

Module description

An implementation of an extractor of a smallest size minimal unsatisfiable subset (smallest MUS, or SMUS)¹²³⁴ and enumerator of SMUSes based on *implicit hitting set enumeration*². This implementation tries to replicate the well-known SMUS extractor Forques². In contrast to Forques, this implementation supports not only plain DIMACS *CNF* formulas but also weighted *WCNF* formulas. As a result, the tool is able to compute and enumerate *optimal* MUSes in case of weighted formulas. On the other hand, this prototype lacks a number of command-line options used in Forques and so it may be less efficient compared to Forques but the performance difference should not be significant.

The file provides a class `OptUx`, which is the basic implementation of the algorithm. It can be applied to any formula in the *CNF* or *WCNF* format.

The implementation can be used as an executable (the list of available command-line options can be shown using `optux.py -h`) in the following way:

```
$ xzcat formula.wcnf.xz
p wcnf 3 6 4
1 1 0
1 2 0
1 3 0
4 -1 -2 0
4 -1 -3 0
4 -2 -3 0

$ optux.py -vvv formula.wcnf.xz
c mcs: 1 2 0
c mcses: 0 unit, 1 disj
c mus: 1 2 0
c cost: 2
c oracle time: 0.0001
```

Alternatively, the algorithm can be accessed and invoked through the standard `import` interface of Python, e.g.

```
>>> from pysat.examples.optux import OptUx
>>> from pysat.formula import WCNF
>>>
```

(continues on next page)

¹ Alexey Ignatiev, Alessandro Previti, Mark H. Liffiton, Joao Marques-Silva. *Smallest MUS Extraction with Minimal Hitting Set Dualization*. CP 2015. pp. 173-182

² Mark H. Liffiton, Maher N. Mneimneh, Ines Lynce, Zaher S. Andraus, Joao Marques-Silva, Karem A. Sakallah. *A branch and bound algorithm for extracting smallest minimal unsatisfiable subformulas*. Constraints An Int. J. 14(4). 2009. pp. 415-442

³ Alexey Ignatiev, Mikolas Janota, Joao Marques-Silva. *Quantified Maximum Satisfiability: A Core-Guided Approach*. SAT 2013. pp. 250-266

⁴ Alexey Ignatiev, Mikolas Janota, Joao Marques-Silva. *Quantified maximum satisfiability*. Constraints An Int. J. 21(2). 2016. pp. 277-302

(continued from previous page)

```
>>> wcnf = WCNF(from_file='formula.wcnf.xz')
>>>
>>> with OptUx(wcnf) as optux:
...     for mus in optux.enumerate():
...         print('mus {0} has cost {1}'.format(mus, optux.cost))
mus [1, 2] has cost 2
mus [1, 3] has cost 2
mus [2, 3] has cost 2
```

As can be seen in the example above, the solver can be instructed either to compute one optimal MUS of an input formula, or to enumerate a given number (or *all*) of its top optimal MUSes.

Module details

class `examples.optux.OptUx`(*formula*, *solver*='g3', *adapt*=False, *cover*=None, *dcalls*=False, *exhaust*=False, *minz*=False, *unsorted*=False, *trim*=False, *verbose*=0)

A simple Python version of the implicit hitting set based optimal MUS extractor and enumerator. Given a (weighted) (partial) CNF formula, i.e. formula in the *WCNF* format, this class can be used to compute a given number of optimal MUS (starting from the *best* one) of the input formula. *OptUx* roughly follows the implementation of Forges[?] but lacks a few additional heuristics, which however aren't applied in Forges by default.

As a result, *OptUx* applies exhaustive *disjoint* minimal correction subset (MCS) enumeration^{?, ?, Page 66, 4} with the incremental use of RC2⁵ as an underlying MaxSAT solver. Once disjoint MCSes are enumerated, they are used to bootstrap a hitting set solver. This implementation uses *Hitman* as a hitting set solver, which is again based on RC2.

Note that in the main implicit hitting enumeration loop of the algorithm, *OptUx* follows Forges in that it does not reduce correction subsets detected to minimal correction subsets. As a result, correction subsets computed in the main loop are added to *Hitman* *unreduced*.

OptUx can use any SAT solver available in PySAT. The default SAT solver to use is g3, which stands for Glucose 3⁶ (see *SolverNames*). Boolean parameters *adapt*, *exhaust*, and *minz* control whether or not the underlying *RC2* oracles should apply detection and adaptation of intrinsic AtMost1 constraints, core exhaustion, and core reduction. Also, unsatisfiable cores can be trimmed if the *trim* parameter is set to a non-zero integer. Finally, verbosity level can be set using the *verbose* parameter.

Two additional optional parameters *unsorted* and *dcalls* can be used to instruct the tool to enumerate MUSes in the unsorted fashion, i.e. optimal MUSes are not guaranteed to go first. For this, *OptUx* applies LBX-like MCS enumeration (it uses *LBX* directly). Parameter *dcalls* can be applied to instruct the underlying MCS enumerator to apply clause D oracle calls.

Finally, one more optional input parameter *cover* is to be used when exhaustive enumeration of MUSes is not necessary and the tool can stop as soon as a given formula is covered by the set of currently computed MUSes. This can be made to work if the soft clauses of *formula* are of size 1.

Parameters

- **formula** (*WCNFPlus*) – (weighted) (partial) CNFPlus formula
- **solver** (*str*) – SAT oracle name
- **adapt** (*bool*) – detect and adapt intrinsic AtMost1 constraints

⁵ Alexey Ignatiev, Antonio Morgado, Joao Marques-Silva. *RC2: an Efficient MaxSAT Solver*. J. Satisf. Boolean Model. Comput. 11(1). 2019. pp. 53-64

⁶ Gilles Audemard, Jean-Marie Lagniez, Laurent Simon. *Improving Glucose for Incremental SAT Solving with Assumptions: Application to MUS Extraction*. SAT 2013. pp. 309-317

- **cover** (*CNFPlus*) – CNFPlus formula to cover when doing MUS enumeration
- **dcalls** (*bool*) – apply clause D oracle calls (for unsorted enumeration only)
- **exhaust** (*bool*) – do core exhaustion
- **minz** (*bool*) – do heuristic core reduction
- **unsorted** (*bool*) – apply unsorted MUS enumeration
- **trim** (*int*) – do core trimming at most this number of times
- **verbose** (*int*) – verbosity level

_disjoint(*formula, solver, adapt, exhaust, minz, trim*)

This method constitutes the preliminary step of the implicit hitting set paradigm of Forges. Namely, it enumerates all the disjoint *minimal correction subsets* (MCSes) of the formula, which will be later used to bootstrap the hitting set solver.

Note that the MaxSAT solver in use is *RC2*. As a result, all the input parameters of the method, namely, *formula*, *solver*, *adapt*, *exhaust*, *minz*, and *trim* - represent the input and the options for the RC2 solver.

Parameters

- **formula** (*WCNF*) – input formula
- **solver** (*str*) – SAT solver name
- **adapt** (*bool*) – detect and adapt AtMost1 constraints
- **exhaust** (*bool*) – exhaust unsatisfiable cores
- **minz** (*bool*) – apply heuristic core minimization
- **trim** (*int*) – trim unsatisfiable cores at most this number of times

_process_soft(*formula*)

The method is for processing the soft clauses of the input formula. Concretely, it checks which soft clauses must be relaxed by a unique selector literal and applies the relaxation.

Parameters **formula** (*WCNF*) – input formula

compute()

This method implements the main look of the implicit hitting set paradigm of Forges to compute a best-cost MUS. The result MUS is returned as a list of integers, each representing a soft clause index.

Return type list(int)

delete()

Explicit destructor of the internal hitting set and SAT oracles.

enumerate()

This is generator method iterating through MUSes and enumerating them until the formula has no more MUSes, or a user decides to stop the process.

Return type list(int)

oracle_time()

This method computes and returns the total SAT solving time involved.

Return type float

1.2.10 RC2 MaxSAT solver (`pysat.examples.rc2`)

List of classes

<code>RC2</code>	Implementation of the basic RC2 algorithm.
<code>RC2Stratified</code>	RC2 augmented with BLO and stratification techniques.

Module description

An implementation of the RC2 algorithm for solving maximum satisfiability. RC2 stands for *relaxable cardinality constraints* (alternatively, *soft cardinality constraints*) and represents an improved version of the OLLITI algorithm, which was described in¹ and² and originally implemented in the [MSCG MaxSAT solver](#).

Initially, this solver was supposed to serve as an example of a possible PySAT usage illustrating how a state-of-the-art MaxSAT algorithm could be implemented in Python and still be efficient. It participated in the [MaxSAT Evaluations 2018](#) and [2019](#) where, surprisingly, it was ranked first in two complete categories: *unweighted* and *weighted*. A brief solver description can be found in³. A more detailed solver description can be found in⁴.

The file implements two classes: `RC2` and `RC2Stratified`. The former class is the basic implementation of the algorithm, which can be applied to a MaxSAT formula in the [WCNFPlus](#) format. The latter class additionally implements Boolean lexicographic optimization (BLO)⁵ and stratification⁶ on top of `RC2`.

The implementation can be used as an executable (the list of available command-line options can be shown using `rc2.py -h`) in the following way:

```
$ xzcat formula.wcnf.xz
p wcnf 3 6 4
1 1 0
1 2 0
1 3 0
4 -1 -2 0
4 -1 -3 0
4 -2 -3 0

$ rc2.py -vv formula.wcnf.xz
c formula: 3 vars, 3 hard, 3 soft
c cost: 1; core sz: 2; soft sz: 2
c cost: 2; core sz: 2; soft sz: 1
s OPTIMUM FOUND
o 2
v -1 -2 3
c oracle time: 0.0001
```

Alternatively, the algorithm can be accessed and invoked through the standard `import` interface of Python, e.g.

¹ António Morgado, Carmine Dodaro, Joao Marques-Silva. *Core-Guided MaxSAT with Soft Cardinality Constraints*. CP 2014. pp. 564-573

² António Morgado, Alexey Ignatiev, Joao Marques-Silva. *MSCG: Robust Core-Guided MaxSAT Solving*. JSAT 9. 2014. pp. 129-134

³ Alexey Ignatiev, António Morgado, Joao Marques-Silva. *RC2: A Python-based MaxSAT Solver*. MaxSAT Evaluation 2018. p. 22

⁴ Alexey Ignatiev, António Morgado, Joao Marques-Silva. *RC2: An Efficient MaxSAT Solver*. MaxSAT Evaluation 2018. JSAT 11. 2019. pp. 53-64

⁵ Joao Marques-Silva, Josep Argelich, Ana Graça, Inês Lynce. *Boolean lexicographic optimization: algorithms & applications*. Ann. Math. Artif. Intell. 62(3-4). 2011. pp. 317-343

⁶ Carlos Ansótegui, Maria Luisa Bonet, Joel Gabàs, Jordi Levy. *Improving WPM2 for (Weighted) Partial MaxSAT*. CP 2013. pp. 117-132

```
>>> from pysat.examples.rc2 import RC2
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF(from_file='formula.wcnf.xz')
>>>
>>> with RC2(wcnf) as rc2:
...     for m in rc2.enumerate():
...         print('model {0} has cost {1}'.format(m, rc2.cost))
model [-1, -2, 3] has cost 2
model [1, -2, -3] has cost 2
model [-1, 2, -3] has cost 2
model [-1, -2, -3] has cost 3
```

As can be seen in the example above, the solver can be instructed either to compute one MaxSAT solution of an input formula, or to enumerate a given number (or *all*) of its top MaxSAT solutions.

Module details

class `examples.rc2.RC2`(*formula*, *solver*='g3', *adapt*=False, *exhaust*=False, *incr*=False, *minz*=False, *trim*=0, *verbose*=0)

Implementation of the basic RC2 algorithm. Given a (weighted) (partial) CNF formula, i.e. formula in the [WCNFPlus](#) format, this class can be used to compute a given number of MaxSAT solutions for the input formula. [RC2](#) roughly follows the implementation of algorithm OLLITI⁷ of MSCG and applies a few heuristics on top of it. These include

- *unsatisfiable core exhaustion* (see method [exhaust_core\(\)](#)),
- *unsatisfiable core reduction* (see method [minimize_core\(\)](#)),
- *intrinsic AtMost1 constraints* (see method [adapt_am1\(\)](#)).

[RC2](#) can use any SAT solver available in PySAT. The default SAT solver to use is g3 (see [SolverNames](#)). Additionally, if Glucose is chosen, the `incr` parameter controls whether to use the incremental mode of Glucose⁷ (turned off by default). Boolean parameters `adapt`, `exhaust`, and `minz` control whether or to apply detection and adaptation of intrinsic AtMost1 constraints, core exhaustion, and core reduction. Unsatisfiable cores can be trimmed if the `trim` parameter is set to a non-zero integer. Finally, verbosity level can be set using the `verbose` parameter.

Parameters

- **formula** ([WCNFPlus](#)) – (weighted) (partial) CNFPlus formula
- **solver** (*str*) – SAT oracle name
- **adapt** (*bool*) – detect and adapt intrinsic AtMost1 constraints
- **exhaust** (*bool*) – do core exhaustion
- **incr** (*bool*) – use incremental mode of Glucose
- **minz** (*bool*) – do heuristic core reduction
- **trim** (*int*) – do core trimming at most this number of times
- **verbose** (*int*) – verbosity level

⁷ Gilles Audemard, Jean-Marie Lagniez, Laurent Simon. *Improving Glucose for Incremental SAT Solving with Assumptions: Application to MUS Extraction*. SAT 2013. pp. 309-317

`_map_extlit()`

Map an external variable to an internal one if necessary.

This method is used when new clauses are added to the formula incrementally, which may result in introducing new variables clashing with the previously used *clause selectors*. The method makes sure no clash occurs, i.e. it maps the original variables used in the new problem clauses to the newly introduced auxiliary variables (see [add_clause\(\)](#)).

Given an integer literal, a fresh literal is returned. The returned integer has the same sign as the input literal.

Parameters `l (int)` – literal to map

Return type `int`

`adapt_am1()`

Detect and adapt intrinsic AtMost1 constraints. Assume there is a subset of soft clauses $\mathcal{S}' \subseteq \mathcal{S}$ s.t. $\sum_{c \in \mathcal{S}'} c \leq 1$, i.e. at most one of the clauses of \mathcal{S}' can be satisfied.

Each AtMost1 relationship between the soft clauses can be detected in the following way. The method traverses all soft clauses of the formula one by one, sets one respective selector literal to true and checks whether some other soft clauses are forced to be false. This is checked by testing if selectors for other soft clauses are unit-propagated to be false. Note that this method for detection of AtMost1 constraints is *incomplete*, because in general unit propagation does not suffice to test whether or not $\mathcal{F} \wedge l_i \models \neg l_j$.

Each intrinsic AtMost1 constraint detected this way is handled by calling [process_am1\(\)](#).

`add_clause(clause, weight=None)`

The method for adding a new hard or soft clause to the problem formula. Although the input formula is to be specified as an argument of the constructor of [RC2](#), adding clauses may be helpful when *enumerating* MaxSAT solutions of the formula. This way, the clauses are added incrementally, i.e. *on the fly*.

The clause to add can be any iterable over integer literals. The additional integer parameter `weight` can be set to meaning the the clause being added is soft having the corresponding weight (note that parameter `weight` is set to `None` by default meaning that the clause is hard).

Also note that besides pure clauses, the method can also expect native cardinality constraints represented as a pair (`lits`, `bound`). Only hard cardinality constraints can be added.

Parameters

- **`clause (iterable(int))`** – a clause to add
- **`weight (int)`** – weight of the clause (if any)

```
>>> from pysat.examples.rc2 import RC2
>>> from pysat.formula import WCNF
>>>
>>> wcnf = WCNF()
>>> wcnf.append([-1, -2]) # adding hard clauses
>>> wcnf.append([-1, -3])
>>>
>>> wcnf.append([1], weight=1) # adding soft clauses
>>> wcnf.append([2], weight=1)
>>> wcnf.append([3], weight=1)
>>>
>>> with RC2(wcnf) as rc2:
...     rc2.compute() # solving the MaxSAT problem
[-1, 2, 3]
...     print(rc2.cost)
```

(continues on next page)

(continued from previous page)

```

1
...     rc2.add_clause([-2, -3]) # adding one more hard clause
...     rc2.compute() # computing another model
[-1, -2, 3]
...     print(rc2.cost)
2

```

compute()

This method can be used for computing one MaxSAT solution, i.e. for computing an assignment satisfying all hard clauses of the input formula and maximizing the sum of weights of satisfied soft clauses. It is a wrapper for the internal `compute_()` method, which does the job, followed by the model extraction.

Note that the method returns `None` if no MaxSAT model exists. The method can be called multiple times, each being followed by blocking the last model. This way one can enumerate top-*k* MaxSAT solutions (this can also be done by calling `enumerate()`).

Returns a MaxSAT model

Return type list(int)

```

>>> from pysat.examples.rc2 import RC2
>>> from pysat.formula import WCNF
>>>
>>> rc2 = RC2(WCNF()) # passing an empty WCNF() formula
>>> rc2.add_clause([-1, -2])
>>> rc2.add_clause([-1, -3])
>>> rc2.add_clause([-2, -3])
>>>
>>> rc2.add_clause([1], weight=1)
>>> rc2.add_clause([2], weight=1)
>>> rc2.add_clause([3], weight=1)
>>>
>>> model = rc2.compute()
>>> print(model)
[-1, -2, 3]
>>> print(rc2.cost)
2
>>> rc2.delete()

```

compute_()

Main core-guided loop, which iteratively calls a SAT oracle, extracts a new unsatisfiable core and processes it. The loop finishes as soon as a satisfiable formula is obtained. If specified in the command line, the method additionally calls `adapt_am1()` to detect and adapt intrinsic AtMost1 constraints before executing the loop.

Return type bool

create_sum(bound=1)

Create a totalizer object encoding a cardinality constraint on the new list of relaxation literals obtained in `process_sels()` and `process_sums()`. The clauses encoding the sum of the relaxation literals are added to the SAT oracle. The sum of the totalizer object is encoded up to the value of the input parameter `bound`, which is set to 1 by default.

Parameters `bound` (int) – right-hand side for the sum to be created

Return type `ITotalizer`

Note that if Minicard is used as a SAT oracle, native cardinality constraints are used instead of *ITotalizer*.

delete()

Explicit destructor of the internal SAT oracle and all the totalizer objects creating during the solving process.

enumerate(block=0)

Enumerate top MaxSAT solutions (from best to worst). The method works as a generator, which iteratively calls *compute()* to compute a MaxSAT model, blocks it internally and returns it.

An optional parameter can be used to enforce computation of MaxSAT models corresponding to different maximal satisfiable subsets (MSSes) or minimal correction subsets (MCSes). To block MSSes, one should set the *block* parameter to 1. To block MCSes, set it to -1. By the default (for blocking MaxSAT models), *block* is set to 0.

Parameters *block* (*int*) – preferred way to block solutions when enumerating

Returns a MaxSAT model

Return type list(int)

```
>>> from pysat.examples.rc2 import RC2
>>> from pysat.formula import WCNF
>>>
>>> rc2 = RC2(WCNF()) # passing an empty WCNF() formula
>>> rc2.add_clause([-1, -2]) # adding clauses "on the fly"
>>> rc2.add_clause([-1, -3])
>>> rc2.add_clause([-2, -3])
>>>
>>> rc2.add_clause([1], weight=1)
>>> rc2.add_clause([2], weight=1)
>>> rc2.add_clause([3], weight=1)
>>>
>>> for model in rc2.enumerate():
...     print(model, rc2.cost)
[-1, -2, 3] 2
[1, -2, -3] 2
[-1, 2, -3] 2
[-1, -2, -3] 3
>>> rc2.delete()
```

exhaust_core(tobj)

Exhaust core by increasing its bound as much as possible. Core exhaustion was originally referred to as *cover optimization* in[?].

Given a totalizer object *tobj* representing a sum of some *relaxation* variables $r \in R$ that augment soft clauses C_r , the idea is to increase the right-hand side of the sum (which is equal to 1 by default) as much as possible, reaching a value k s.t. formula $\mathcal{H} \wedge C_r \wedge (\sum_{r \in R} r \leq k)$ is still unsatisfiable while increasing it further makes the formula satisfiable (here \mathcal{H} denotes the hard part of the formula).

The rationale is that calling an oracle incrementally on a series of slightly modified formulas focusing only on the recently computed unsatisfiable core and disregarding the rest of the formula may be practically effective.

filter_assumps()

Filter out unnecessary selectors and sums from the list of assumption literals. The corresponding values are also removed from the dictionaries of bounds and weights.

Note that assumptions marked as garbage are collected in the core processing methods, i.e. in `process_core()`, `process_sels()`, and `process_sums()`.

`get_core()`

Extract unsatisfiable core. The result of the procedure is stored in variable `self.core`. If necessary, core trimming and also heuristic core reduction is applied depending on the command-line options. A *minimum weight* of the core is computed and stored in `self.minw`. Finally, the core is divided into two parts:

1. clause selectors (`self.core_sels`),
2. sum assumptions (`self.core_sums`).

`init(formula, incr=False)`

Initialize the internal SAT oracle. The oracle is used incrementally and so it is initialized only once when constructing an object of class `RC2`. Given an input *WCNFPlus* formula, the method bootstraps the oracle with its hard clauses. It also augments the soft clauses with “fresh” selectors and adds them to the oracle afterwards.

Optional input parameter `incr` (False by default) regulates whether or not Glucose’s incremental mode^{Page 70, 7} is turned on.

Parameters

- **formula** (*WCNFPlus*) – input formula
- **incr** (*bool*) – apply incremental mode of Glucose

`minimize_core()`

Reduce a previously extracted core and compute an over-approximation of an MUS. This is done using the simple deletion-based MUS extraction algorithm.

The idea is to try to deactivate soft clauses of the unsatisfiable core one by one while checking if the remaining soft clauses together with the hard part of the formula are unsatisfiable. Clauses that are necessary for preserving unsatisfiability comprise an MUS of the input formula (it is contained in the given unsatisfiable core) and are reported as a result of the procedure.

During this core minimization procedure, all SAT calls are dropped after obtaining 1000 conflicts.

`oracle_time()`

Report the total SAT solving time.

`process_am1(am1)`

Process an AtMost1 relation detected by `adapt_am1()`. Note that given a set of soft clauses S' at most one of which can be satisfied, one can immediately conclude that the formula has cost at least $|S'| - 1$ (assuming *unweighted* MaxSAT). Furthermore, it is safe to replace all clauses of S' with a single soft clause $\sum_{c \in S'} c$.

Here, input parameter `am1` plays the role of subset S' mentioned above. The procedure bumps the MaxSAT cost by `self.minw * (len(am1) - 1)`.

All soft clauses involved in `am1` are replaced by a single soft clause, which is a disjunction of the selectors of clauses in `am1`. The weight of the new soft clause is set to `self.minw`.

Parameters `am1` (`list(int)`) – a list of selectors connected by an AtMost1 constraint

`process_core()`

The method deals with a core found previously in `get_core()`. Clause selectors `self.core_sels` and sum assumptions involved in the core are treated separately of each other. This is handled by calling methods `process_sels()` and `process_sums()`, respectively. Whenever necessary, both methods relax the core literals, which is followed by creating a new totalizer object encoding the sum of the new relaxation variables. The totalizer object can be “exhausted” depending on the option.

process_sels()

Process soft clause selectors participating in a new core. The negation $\neg s$ of each selector literal s participating in the unsatisfiable core is added to the list of relaxation literals, which will be later used to create a new totalizer object in `create_sum()`.

If the weight associated with a selector is equal to the minimal weight of the core, e.g. `self.minw`, the selector is marked as garbage and will be removed in `filter_assumps()`. Otherwise, the clause is split as described in[?].

process_sums()

Process cardinality sums participating in a new core. Whenever necessary, some of the sum assumptions are removed or split (depending on the value of `self.minw`). Deleted sums are marked as garbage and are dealt with in `filter_assumps()`.

In some cases, the process involves updating the right-hand sides of the existing cardinality sums (see the call to `update_sum()`). The overall procedure is detailed in[?].

set_bound(*tobj*, *rhs*)

Given a totalizer sum and its right-hand side to be enforced, the method creates a new sum assumption literal, which will be used in the following SAT oracle calls.

Parameters

- **tobj** (*ITotalizer*) – totalizer sum
- **rhs** (*int*) – right-hand side

trim_core()

This method trims a previously extracted unsatisfiable core at most a given number of times. If a fixed point is reached before that, the method returns.

update_sum(*assump*)

The method is used to increase the bound for a given totalizer sum. The totalizer object is identified by the input parameter `assump`, which is an assumption literal associated with the totalizer object.

The method increases the bound for the totalizer sum, which involves adding the corresponding new clauses to the internal SAT oracle.

The method returns the totalizer object followed by the new bound obtained.

Parameters `assump` (*int*) – assumption literal associated with the sum

Return type *ITotalizer*, *int*

Note that if Minicard is used as a SAT oracle, native cardinality constraints are used instead of *ITotalizer*.

```
class examples.rc2.RC2Stratified(formula, solver='g3', adapt=False, blo='div', exhaust=False, incr=False,
                               minz=False, nohard=False, trim=0, verbose=0)
```

RC2 augmented with BLO and stratification techniques. Although class `RC2` can deal with weighted formulas, there are situations when it is necessary to apply additional heuristics to improve the performance of the solver on weighted MaxSAT formulas. This class extends capabilities of `RC2` with two heuristics, namely

1. Boolean lexicographic optimization (BLO)[?]
2. diversity-based stratification[?]
3. cluster-based stratification

To specify which heuristics to apply, a user can assign the `blo` parameter to one of the values (by default it is set to 'div'):

- 'basic' ('BLO' only)

- `div` ('BLO' + diversity-based stratification)
- `cluster` ('BLO' + cluster-based stratification)
- `full` ('BLO' + diversity- + cluster-based stratification)

Except for the aforementioned additional techniques, every other component of the solver remains as in the base class `RC2`. Therefore, a user is referred to the documentation of `RC2` for details.

`activate_clauses(beg)`

This method is used for activating the clauses that belong to optimization levels up to the newly computed level. It also reactivates previously deactivated clauses (see `process_sels()` and `process_sums()` for details).

`compute()`

This method solves the MaxSAT problem iteratively. Each optimization level is tackled the standard way, i.e. by calling `compute_()`. A new level is started by calling `next_level()` and finished by calling `finish_level()`. Each new optimization level activates more soft clauses by invoking `activate_clauses()`.

`finish_level()`

This method does postprocessing of the current optimization level after it is solved. This includes *hardening* some of the soft clauses (depending on their remaining weights) and also garbage collection.

`init_wstr()`

Compute and initialize optimization levels for BLO and stratification. This method is invoked once, from the constructor of an object of `RC2Stratified`. Given the weights of the soft clauses, the method divides the MaxSAT problem into several optimization levels.

`next_level()`

Compute the next optimization level (starting from the current one). The procedure represents a loop, each iteration of which checks whether or not one of the conditions holds:

- partial BLO condition
- diversity-based stratification condition
- cluster-based stratification condition

If any of these holds, the loop stops.

`process_am1(am1)`

Due to the solving process involving multiple optimization levels to be treated individually, new soft clauses for the detected intrinsic AtMost1 constraints should be remembered. The method is a slightly modified version of the base method `RC2.process_am1()` taking care of this.

`process_sels()`

A redefined version of `RC2.process_sels()`. The only modification affects the clauses whose weight after splitting becomes less than the weight of the current optimization level. Such clauses are deactivated and to be reactivated at a later stage.

`process_sums()`

A redefined version of `RC2.process_sums()`. The only modification affects the clauses whose weight after splitting becomes less than the weight of the current optimization level. Such clauses are deactivated and to be reactivated at a later stage.

PYTHON MODULE INDEX

e

- `examples.fm`, 44
- `examples.genhard`, 46
- `examples.hitman`, 49
- `examples.lbx`, 54
- `examples.lsu`, 57
- `examples.mcs1s`, 60
- `examples.models`, 63
- `examples.musx`, 64
- `examples.optux`, 66
- `examples.rc2`, 69

p

- `pysat.card`, 3
- `pysat.formula`, 9
- `pysat.pb`, 29
- `pysat.solvers`, 31

Symbols

`_assert_lt()` (*examples.lsu.LSU method*), 58
`_assert_lt()` (*examples.lsu.LSUPlus method*), 59
`_compute()` (*examples.fm.FM method*), 45
`_compute()` (*examples.lbx.LBX method*), 55
`_compute()` (*examples.mcsls.MCSls method*), 61
`_compute()` (*examples.musx.MUSX method*), 65
`_disjoint()` (*examples.optux.OptUx method*), 68
`_filter_satisfied()` (*examples.lbx.LBX method*), 55
`_get_model_cost()` (*examples.lsu.LSU method*), 58
`_init()` (*examples.lsu.LSU method*), 59
`_map_extlit()` (*examples.lbx.LBX method*), 55
`_map_extlit()` (*examples.mcsls.MCSls method*), 61
`_map_extlit()` (*examples.rc2.RC2 method*), 70
`_overapprox()` (*examples.mcsls.MCSls method*), 61
`_process_soft()` (*examples.optux.OptUx method*), 68
`_satisfied()` (*examples.lbx.LBX method*), 55

A

`accum_stats()` (*pysat.solvers.Solver method*), 35
`activate_clauses()` (*examples.rc2.RC2Stratified method*), 76
`adapt_am1()` (*examples.rc2.RC2 method*), 71
`add_atmost()` (*pysat.solvers.Solver method*), 35
`add_clause()` (*examples.lbx.LBX method*), 56
`add_clause()` (*examples.mcsls.MCSls method*), 61
`add_clause()` (*examples.rc2.RC2 method*), 71
`add_clause()` (*pysat.solvers.Solver method*), 35
`add_hard()` (*examples.hitman.Hitman method*), 52
`append()` (*pysat.formula.CNF method*), 12
`append()` (*pysat.formula.CNFPlus method*), 17
`append()` (*pysat.formula.WCNF method*), 21
`append()` (*pysat.formula.WCNFPlus method*), 26
`append_formula()` (*pysat.solvers.Solver method*), 36
`atleast()` (*pysat.card.CardEnc class method*), 4
`atleast()` (*pysat.pb.PBEnc class method*), 30
`atmost()` (*pysat.card.CardEnc class method*), 5
`atmost()` (*pysat.pb.PBEnc class method*), 30
`Atom` (*class in examples.hitman*), 51

B

`block()` (*examples.hitman.Hitman method*), 52

`block()` (*examples.lbx.LBX method*), 56
`block()` (*examples.mcsls.MCSls method*), 62

C

`CardEnc` (*class in pysat.card*), 4
`CB` (*class in examples.genhard*), 48
`clear_interrupt()` (*examples.lsu.LSU method*), 59
`clear_interrupt()` (*pysat.solvers.Solver method*), 36
`CNF` (*class in pysat.formula*), 11
`CNFPlus` (*class in pysat.formula*), 16
`compute()` (*examples.fm.FM method*), 45
`compute()` (*examples.lbx.LBX method*), 56
`compute()` (*examples.mcsls.MCSls method*), 62
`compute()` (*examples.musx.MUSX method*), 65
`compute()` (*examples.optux.OptUx method*), 68
`compute()` (*examples.rc2.RC2 method*), 72
`compute()` (*examples.rc2.RC2Stratified method*), 76
`compute_()` (*examples.rc2.RC2 method*), 72
`conf_budget()` (*pysat.solvers.Solver method*), 36
`copy()` (*pysat.formula.CNF method*), 12
`copy()` (*pysat.formula.CNFPlus method*), 17
`copy()` (*pysat.formula.WCNF method*), 22
`copy()` (*pysat.formula.WCNFPlus method*), 27
`create_sum()` (*examples.rc2.RC2 method*), 72

D

`delete()` (*examples.fm.FM method*), 45
`delete()` (*examples.hitman.Hitman method*), 53
`delete()` (*examples.lbx.LBX method*), 56
`delete()` (*examples.lsu.LSU method*), 59
`delete()` (*examples.mcsls.MCSls method*), 62
`delete()` (*examples.musx.MUSX method*), 65
`delete()` (*examples.optux.OptUx method*), 68
`delete()` (*examples.rc2.RC2 method*), 73
`delete()` (*pysat.card.ITotalizer method*), 6
`delete()` (*pysat.solvers.Solver method*), 36
`do_cld_check()` (*examples.lbx.LBX method*), 56
`do_cld_check()` (*examples.mcsls.MCSls method*), 62

E

`EncType` (*class in pysat.card*), 5
`EncType` (*class in pysat.pb*), 30

`enum_models()` (*pysat.solvers.Solver method*), 36
`enumerate()` (*examples.hitman.Hitman method*), 53
`enumerate()` (*examples.lbx.LBX method*), 57
`enumerate()` (*examples.mcsIs.MCSIs method*), 62
`enumerate()` (*examples.optux.OptUx method*), 68
`enumerate()` (*examples.rc2.RC2 method*), 73
`enumerate_models()` (*in module examples.models*), 63
`equals()` (*pysat.card.CardEnc class method*), 5
`equals()` (*pysat.pb.PBEnc class method*), 30
`examples.fm`
 module, 44
`examples.genhard`
 module, 46
`examples.hitman`
 module, 49
`examples.lbx`
 module, 54
`examples.lsu`
 module, 57
`examples.mcsIs`
 module, 60
`examples.models`
 module, 63
`examples.musx`
 module, 64
`examples.optux`
 module, 66
`examples.rc2`
 module, 69
`exhaust_core()` (*examples.rc2.RC2 method*), 73
`extend()` (*pysat.card.ITotalizer method*), 6
`extend()` (*pysat.formula.CNF method*), 12
`extend()` (*pysat.formula.CNFPlus method*), 18
`extend()` (*pysat.formula.WCNF method*), 22

F

`filter_assumps()` (*examples.rc2.RC2 method*), 73
`finish_level()` (*examples.rc2.RC2Stratified method*), 76
`FM` (*class in examples.fm*), 45
`found_optimum()` (*examples.lsu.LSU method*), 59
`from_aiger()` (*pysat.formula.CNF method*), 12
`from_clauses()` (*pysat.formula.CNF method*), 13
`from_file()` (*pysat.formula.CNF method*), 13
`from_file()` (*pysat.formula.WCNF method*), 22
`from_fp()` (*pysat.formula.CNF method*), 14
`from_fp()` (*pysat.formula.CNFPlus method*), 18
`from_fp()` (*pysat.formula.WCNF method*), 23
`from_fp()` (*pysat.formula.WCNFPlus method*), 27
`from_string()` (*pysat.formula.CNF method*), 14
`from_string()` (*pysat.formula.WCNF method*), 23

G

`geq()` (*pysat.pb.PBEnc class method*), 31

`get()` (*examples.hitman.Hitman method*), 53
`get_core()` (*examples.rc2.RC2 method*), 74
`get_core()` (*pysat.solvers.Solver method*), 37
`get_model()` (*examples.lsu.LSU method*), 59
`get_model()` (*pysat.solvers.Solver method*), 37
`get_proof()` (*pysat.solvers.Solver method*), 38
`get_status()` (*pysat.solvers.Solver method*), 38
`GT` (*class in examples.genhard*), 48

H

`hit()` (*examples.hitman.Hitman method*), 53
`Hitman` (*class in examples.hitman*), 51

I

`id()` (*pysat.formula.IDPool method*), 20
`IDPool` (*class in pysat.formula*), 20
`increase()` (*pysat.card.ITotalizer method*), 7
`init()` (*examples.fm.FM method*), 45
`init()` (*examples.hitman.Hitman method*), 53
`init()` (*examples.rc2.RC2 method*), 74
`init_wstr()` (*examples.rc2.RC2Stratified method*), 76
`interrupt()` (*examples.lsu.LSU method*), 59
`interrupt()` (*pysat.solvers.Solver method*), 38
`ITotalizer` (*class in pysat.card*), 5

L

`LBX` (*class in examples.lbx*), 55
`leq()` (*pysat.pb.PBEnc class method*), 31
`LSU` (*class in examples.lsu*), 58
`LSUPlus` (*class in examples.lsu*), 59

M

`MCSIs` (*class in examples.mcsIs*), 61
`merge_with()` (*pysat.card.ITotalizer method*), 7
`minimize_core()` (*examples.rc2.RC2 method*), 74
`module`

examples.fm, 44
 examples.genhard, 46
 examples.hitman, 49
 examples.lbx, 54
 examples.lsu, 57
 examples.mcsIs, 60
 examples.models, 63
 examples.musx, 64
 examples.optux, 66
 examples.rc2, 69
 pysat.card, 3
 pysat.formula, 9
 pysat.pb, 29
 pysat.solvers, 31

`MUSX` (*class in examples.musx*), 65

N

`negate()` (*pysat.formula.CNF method*), 14

`new()` (*pysat.card.ITotalizer method*), 8
`new()` (*pysat.solvers.Solver method*), 39
`next_level()` (*examples.rc2.RC2Stratified method*), 76
`nof_clauses()` (*pysat.solvers.Solver method*), 39
`nof_vars()` (*pysat.solvers.Solver method*), 39
`normalize_negatives()` (*pysat.formula.WCNF method*), 24
`NoSuchEncodingError`, 8, 30
`NoSuchSolverError`, 33

O

`obj()` (*pysat.formula.IDPool method*), 21
`occupy()` (*pysat.formula.IDPool method*), 21
`OptUx` (*class in examples.optux*), 67
`oracle_time()` (*examples.fm.FM method*), 45
`oracle_time()` (*examples.hitman.Hitman method*), 53
`oracle_time()` (*examples.lbx.LBX method*), 57
`oracle_time()` (*examples.lsu.LSU method*), 59
`oracle_time()` (*examples.mcsls.MCSls method*), 62
`oracle_time()` (*examples.musx.MUSX method*), 65
`oracle_time()` (*examples.optux.OptUx method*), 68
`oracle_time()` (*examples.rc2.RC2 method*), 74

P

`PAR` (*class in examples.genhard*), 49
`PBEnc` (*class in pysat.pb*), 30
`PHP` (*class in examples.genhard*), 49
`process_am1()` (*examples.rc2.RC2 method*), 74
`process_am1()` (*examples.rc2.RC2Stratified method*), 76
`process_core()` (*examples.rc2.RC2 method*), 74
`process_sels()` (*examples.rc2.RC2 method*), 74
`process_sels()` (*examples.rc2.RC2Stratified method*), 76
`process_sums()` (*examples.rc2.RC2 method*), 75
`process_sums()` (*examples.rc2.RC2Stratified method*), 76
`prop_budget()` (*pysat.solvers.Solver method*), 39
`propagate()` (*pysat.solvers.Solver method*), 40
`pysat.card`
 module, 3
`pysat.formula`
 module, 9
`pysat.pb`
 module, 29
`pysat.solvers`
 module, 31

R

`RC2` (*class in examples.rc2*), 70
`RC2Stratified` (*class in examples.rc2*), 75
`reinit()` (*examples.fm.FM method*), 45
`relax_core()` (*examples.fm.FM method*), 46
`remove_unit_core()` (*examples.fm.FM method*), 46

`restart()` (*pysat.formula.IDPool method*), 21

S

`set_bound()` (*examples.rc2.RC2 method*), 75
`set_phases()` (*pysat.solvers.Solver method*), 40
`solve()` (*examples.lsu.LSU method*), 59
`solve()` (*pysat.solvers.Solver method*), 41
`solve_limited()` (*pysat.solvers.Solver method*), 41
`Solver` (*class in pysat.solvers*), 33
`SolverNames` (*class in pysat.solvers*), 43
`split_core()` (*examples.fm.FM method*), 46
`supports_atmost()` (*pysat.solvers.Solver method*), 42

T

`time()` (*pysat.solvers.Solver method*), 42
`time_accum()` (*pysat.solvers.Solver method*), 42
`to_alien()` (*pysat.formula.CNF method*), 15
`to_alien()` (*pysat.formula.CNFPlus method*), 18
`to_alien()` (*pysat.formula.WCNF method*), 24
`to_alien()` (*pysat.formula.WCNFPlus method*), 27
`to_file()` (*pysat.formula.CNF method*), 15
`to_file()` (*pysat.formula.WCNF method*), 24
`to_fp()` (*pysat.formula.CNF method*), 16
`to_fp()` (*pysat.formula.CNFPlus method*), 19
`to_fp()` (*pysat.formula.WCNF method*), 25
`to_fp()` (*pysat.formula.WCNFPlus method*), 28
`treat_core()` (*examples.fm.FM method*), 46
`trim_core()` (*examples.rc2.RC2 method*), 75

U

`UnsupportedBound`, 8
`unweighted()` (*pysat.formula.WCNF method*), 25
`unweighted()` (*pysat.formula.WCNFPlus method*), 28
`update_sum()` (*examples.rc2.RC2 method*), 75

W

`WCNF` (*class in pysat.formula*), 21
`WCNFPlus` (*class in pysat.formula*), 25
`weighted()` (*pysat.formula.CNF method*), 16
`weighted()` (*pysat.formula.CNFPlus method*), 19
`with_traceback()` (*pysat.card.NoSuchEncodingError method*), 8
`with_traceback()` (*pysat.card.UnsupportedBound method*), 8
`with_traceback()` (*pysat.pb.NoSuchEncodingError method*), 30
`with_traceback()` (*pysat.solvers.NoSuchSolverError method*), 33